CS10 – Computer Architecture and Organization

My Notes

Leonard Mohr

Student at Foothill College

leonardmohr@gmail.com

Contents

Computer Abstractions and Technology

Introduction

Decimal term	Abbreviation	Value	Binary term	Abbreviation	Value	% Larger
kilobyte	KB	10^3	kibibyte	KiB	2^{10}	2%
megabyte	MB	10^6	mebibyte	MiB	2^{20}	5%
gigabyte	GB	10^9	gibibyte	GiB	2^{30}	7%
terabyte	TB	10^{12}	tebibyte	TiB	2^{40}	10%
petabyte	PB	10^{15}	pebibyte	PiB	2^{50}	13%
exabyte	EB	10^{18}	exbibyte	EiB	2^{60}	15%
zettabyte	ZB	10^{21}	zebibyte	ZiB	2^{70}	18%
yottabyte	YB	10^{24}	yobibyte	YiB	2^{80}	21%

Figure 1. We can describe storage in binary or decimal notation. We are much more familiar with the decimal term. Note also the size difference between the two: binary term gets progressivly larger.

1.1.1   Program Performance

One of the main goals for both the hardware designer and the software designer is to improve performance. This can be achieved in different ways (just think about how quick sort is much faster than bubble sort, and the apple M1 processor is much faster than the intel processors they replaced).

Hardware or software component	How this component affects performance	Where is this topic covered?
Algorithm	Determines both the number of source-level statements and the number of I/O operations executed	Other books!
Programming language, compiler, and architecture	Determines the number of computer instructions for each source-level statement	Chapters 2 and 3
Processor and memory system	Determines how fast instructions can be executed	Chapters 4, 5, and 6
I/O system (hardware and operating system)	Determines how fast I/O operations may be executed	Chapters 4, 5, and 6

Figure 2. Here we can see the various ways program performance can be improved.

Check Yourself

1. As mentioned earlier, both the software and hardware affect the performance of a program. Can you think of examples where each of the following is the right place to look for a performance bottleneck?

   a) The algorithm chosen

   > If a program that sorts a really long list of names is taking a long time, you might want to look at the algorithm being used.

   b) The programming language or compiler

> If your program could is not language dependent, you could look at using a compiled language (like C), as opposed to an interpreted language like Python.

c) The operating system

> If one program runs well, but two at a time don't, then maybe the operating system isn't distributing it's resources efficiently.

d) The processor

> If the computer in general is using a lot of energy, you might want to look at the processor. Similarily, if you are doing compute heavy tasks such as vdeo editing, you might need a processor upgrade.

e) The I/O system and devices

> If it takes a long time to write to a hard drive, you might look at upgrading to a SSD.

Subsection 1.2
Eight Great Ideas in Computer Architecture

Figure 3. Moore's Law in action.

1. **Moore's Law**
   Moore's Law resulted from a 1965 prediction that integrated circuit (IC) resources would double every 18-24 months.

2. **Use Abstraction to Simplify Design**
   To increase productivity, by design both computer architects and programmers try use abstractions to hide the lower-level details and provide a simpler to work with. For example, the operating system abstracts away the complexity of the memory system so that programs are provided with a much simplified view of the memory, and the details are handled by the operating system.

3. **Make the Common Case Fast**
   Better performance gains can be made if you optimize for what the prgram is going to do most often.

4. **Performance via Parallelism**
   We can make a program a lot faster by performing multiple tasks at once. This is especially true with the advent of multi-core processors.

5. **Performance via Pipelining**
   If you are moving a lot of bricks from one place to another (using just manpower) it would be a lot more efficient to set up a line of people, and pass the bricks down the line, than to just have everyone running back and forth. The same principle can be used in computers.

6. **Performance via Prediction**
   When a processor encounters an if statement, it might say that on average the result is true, and procede like it is, so that it keep going, and then the if qualifier can be processed at a later date.

2

7. **Hierarchy of Memories**
   By using different types of memory, from really small and really fast, to really large and slow, a memory hierarchy is created. Caches give the programmer the illusion that main memory is nearly as fast as the top of the hierarchy and nearly as big and cheap as the bottom of the hierarchy.

8. **Dependability via Redundancy**
   Because we are sad when a computer dies, we want to prevent the death of the computer. To do this we make them dependable. This can be achieved by including redundant components that can both take over in the event of a failure as well as detect when a failure has occured.

SUBSECTION 1.3
Below Your Program

The Operating System is a great example of abstraction. Some of its most important functions are:

- Handling basic input and output operations

- Allocating storage and memory

- Provided for protected sharing of the computer among multiple applications using it simultaneously.

Another example of abstraction is high-level programming languages like C. When computers first came about, programmers programmed in binary; they wrote their programs in 1's and 0's (the language of the computer). Since that was tedious, they invented the assembler which would convert assembly language into binary code. Next came the compiler which would convert higher order languages into assembly.

Another benefit of the programming languages is it allows one to chose the specific language that is best for the task. Also, programs don't have to be written for a specific processor, since the compiler and assembler can package it for different computers.

SUBSECTION 1.4
Under the Covers

The five classic components of a computer are input, output, memory, datapath, and control, with the last two sometimes combined and called the processor. This organization is independent of hardware technology: you can place every piece of every computer, past and present, into one of these five categories.

1.4.1   Parts of the Computer

- Integrated Circuit: Also called a chip. A device combining dozens to millions of transistors.

- Central Processing Unit (CPU): Also called the processor. The active part of the computer, which contains the datapath and control and which adds numbers, tests numbers, signals I/O devices to activate, and so on.

- Datapath: The component of the processor that performs arithmetic operations

- Control: The component of the processor that commands the datapath, memory, and I/O devices according to the instructions of the program.

- Memory: The storage area in which programs are kept when they are running and that contains the data needed by the running programs.


Figure 4. Here we see the flow of information in a computer. The processor gets instructions and data from memory. Input writes data to memory, and output reads data from memory. Control sends the signals that determine the operations of the datapath, memory, input, and output.

- Dynamic Random Access Memory (DRAM): Memory built as an integrated circuit; it provides random access to any location Access times are 50 nanoseconds and cost per gigabyte in 2012 was $5 to $10.

- Cache Memory: Consists of a small, fast memory that acts as a buffer for the DRAM memory.

- Static Random Access Memory: SRAM is faster but less dense, and hence more expensive, than DRAM.

- Instruction Set Architector: The interface between the hardware and the lowest-level software. This includes all the information necessary to write a machine language program that will run correctly, including instructions, registers, memory access, I/O, and so on.

- Application Binary Interface (ABI): The user portion of the instruction set plus the operating system interfaces used by application programmers. It defines a standard for binary portability across computers.

Performance

When talking about computers, we often want to look at the performance of the computer. But how can we define performance?

> **Definition 1**   **Response Time**
> Also referred to as **execution time**. This is the total time required for the computer to complete a task, including disk access, memory aaccesses, I/O activities, operating system overhead, CPU execution time, and so on.

> **Definition 2**   **Throughput / Bandwidth**
> This measures the number of tasks computed per unit time.

For the time being, we will mostly be looking at response time. So that a higher number represents a machine with better performance we will say that

$$\text{Performance }_X = \frac{1}{\text{Execution time }_X}.$$

Also, we often want to say that computer "X is n times as fast as Y", to compute n we say:

$$\frac{\text{Performance }_X}{\text{Performance }_Y} = n.$$

CPU execution time or simply **CPU time** is the actual time the CPU spends computing for a specific task.

1.6.1 Clock Rate and Clock Period

Designers refer to the length of a **clock period** both as the time for a complete clock cycle (e.g., 250 picoseconds) and as the *clock rate* (e.g., 4 gigahertz, or 4GHz), which is the inverse of the clock period.

For example a clock rate of

$$
\begin{aligned}
4\text{GHz} &= 4,000,000,000 \frac{\text{cycles}}{\text{sec}} \\
&\equiv \frac{1 \text{ cycle}}{4,000,000,000 \text{ second}} \\
&= 0.00000000025 \frac{\text{cycle}}{\text{second}} \\
&= \frac{1}{250} \frac{\text{cycle}}{\text{picosecond}}.
\end{aligned}
$$

We can relate clock cycles and clock cycle time to CPU time:

$$
\begin{aligned}
\text{CPU execution time} &= \text{CPU clock cycles} \\
\text{for a program} & \quad \text{for a program} \times \text{ Clock cycle time}
\end{aligned}
$$

or alternatively,

$$
\begin{aligned}
\text{CPU execution time} &= \frac{\text{CPU clock cycles for a program}}{\text{Clock rate.}} \\
\text{for a program}
\end{aligned}
$$

1.6.2   Instruction Performance

In addition to thinking about clock cycles and CPU time, we also need to think about number of instructions there are in a program and the time each instruction takes:

$$
\text{CPU clock cycles} = \text{Instructions for a Program} \times \frac{\text{average clock cycles}}{\text{instruction}}
$$

Clock Cycles per Instruction (CPI) is the average number of clock cycles per instruction for a program or program fragement.

1.6.3   The Classic CPU Performance Equation

Putting everything from above together we have

$$
\text{CPU time} = \text{Instruction count} \times \text{CPI} \times \text{ Clock cycle time}
$$

and

$$
\text{CPU time} = \frac{\text{Instruction count} \times \text{CPI}}{\text{Clock rate}}.
$$

We also have

$$
\text{Time} = \text{Seconds / Program} = \frac{\text{Instructions}}{\text{Program}} \times \frac{\text{Clock cycles}}{\text{Instruction}} \times \frac{\text{Seconds}}{\text{Clock cycle.}}
$$

Processors these days can vary their clock rates. For example Intel Core i7 chips temporarily increase clock rate by 10% until the chip gets too warm. Thus we need to use the average clock rate for a program.

Figure 5.

Fallacies and Pitfalls

1.10.1   Amdahl's Law

Definition 3 | **Amdahl's Law**
A rule stating that the performance enhancement possible with a given improvement

is limited by the amount that the improved feature is used:

$$\text{Execution time after improvement} = \frac{\text{Execution time affected by improvement}}{\text{Amount of improvement}}$$
$$+ \;\text{Execution time unaffected}$$

where Amount of improvement $= n$.

We thus see that there is only so much benefit that can be achieved by improving a part of the program that is rarely used. For this reason we want to make the common case fast!

1.10.2 MIPS

One should be wary about using only a subset of the performance equation as a performance metric; one can't determine performance just by looking at clock rate, instruction count, or CPI alone.

An alternative to time is **MIPS (million instructions per second)**:

$$\text{MIPS} = \frac{\text{Instruction count}}{\text{Execution time } \times 10^6}.$$

There are a couple problems with MIPS:

- MIPS specifies the instruction execution rate but does not take into account the capabilities of these instructions. Therefore we can't compare computers with different instruction sets.

- MIPS varies between programs on the same computer; thus a computer cannot have a single MIPS rating.

$$\text{MIPS} = \frac{\text{Instruction count}}{\frac{\text{Instruction count } \times \text{ CPI}}{\text{Clock rate}} \times 10^6} = \frac{\text{Clock rate}}{\text{CPI} \times 10^6}.$$

1.10.3 Check Yourself

Consider the following performance measurements for a program:

Measurement	Computer A	Computer B
Instruction count	10 billion	8 billion
Clock rate	4 GHz	4 GHz
CPI	1.0	1.1

Which computer is faster and which has the higher MIPS rating?

First looking at computer A.

$$\begin{aligned}
\text{Time}(A) &= \frac{\text{CPU Clock Cycles}}{\text{Clock Rate}} \\
&= \frac{\text{Instruction Count } \times \text{ CPI}}{\text{Clock Rate}} \\
&= \frac{10 \times 10^9 \text{ instructions} \times \frac{1 \text{ cycle}}{\text{instruction}}}{4 \times 10^9 \frac{\text{cycles}}{\text{second}}} \\
&= 2.5 \text{ seconds.}
\end{aligned}$$

$$\begin{aligned}
\text{MIPS}(A) &= \frac{\text{Instruction Count}}{\text{Execution Time } \times 10^6} \\
&= \frac{10 \times 10^9 \text{ instructions}}{2.5 \text{ seconds } \times 10^6} \\
&= \frac{4 \times 10^3 \text{ million instructions}}{\text{second}}.
\end{aligned}$$

Now looking at computer B.

$$\begin{aligned}
\text{Time}(B) &= \frac{\text{CPU Clock Cycles}}{\text{Clock Rate}} \\
&= \frac{\text{Instruction Count } \times \text{ CPI}}{\text{Clock Rate}} \\
&= \frac{8 \times 10^9 \text{ instructions} \times \frac{1.1 \text{ cycles}}{\text{instruction}}}{4 \times 10^9 \frac{\text{cycles}}{\text{second}}} \\
&= 2.2 \text{ seconds}.
\end{aligned}$$

$$\begin{aligned}
\text{MIPS}(B) &= \frac{\text{Instruction Count}}{\text{Execution Time } \times 10^6} \\
&= \frac{8 \times 10^9 \text{ instructions}}{2.2 \text{ seconds } \times 10^6} \\
&\approx \frac{3.6 \times 10^3 \text{ million instructions}}{\text{second}}.
\end{aligned}$$

We thus see that Computer A has a higher MIPS score but runs slower.

Instructions: Language of the Computer

Operations of the Computer Hardware

2.2.1 MIPS Operands

Name	Example	Comments
32 registers	$s0-$s7, $t0-$t9, $zero, $a0-$a3, $v0-$v1, $gp, $fp, $sp, $ra, $at	Fast locations for data. In MIPS, data must be in registers to perform arithmetic, register $zero always equals 0, and register $at is reserved by the assembler to handle large constants.
2^{30} memory words	Memory[0], Memory[4], . . . , Memory[4294967292]	Accessed only by data transfer instructions. MIPS uses byte addresses, so sequential word addresses differ by 4. Memory holds data structures, arrays, and spilled registers.

2.2.2 MIPS Instructions

Category	Instruction	Example	Meaning	Comments
Arithmetic	add	`add $s1,$s2,$s3`	$s1 = $s2 + $s3	Three register operands
	subtract	`sub $s1,$s2,$s3`	$s1 = $s2 - $s3	Three register operands
	add immediate	`addi $s1,$s2,20`	$s1 = $s2 + 20	Used to add constants
Data transfer	load word	`lw $s1,20($s2)`	$s1 = Memory[$s2 + 20]	Word from memory to register
	store word	`sw $s1,20($s2)`	Memory[$s2 + 20] = $s1	Word from register to memory
	load half	`lh $s1,20($s2)`	$s1 = Memory[$s2 + 20]	Halfword memory to register
	load half unsigned	`lhu $s1,20($s2)`	$s1 = Memory[$s2 + 20]	Halfword memory to register
	store half	`sh $s1,20($s2)`	Memory[$s2 + 20] = $s1	Halfword register to memory
	load byte	`lb $s1,20($s2)`	$s1 = Memory[$s2 + 20]	Byte from memory to register
	load byte unsigned	`lbu $s1,20($s2)`	$s1 = Memory[$s2 + 20]	Byte from memory to register
	store byte	`sb $s1,20($s2)`	Memory[$s2 + 20] = $s1	Byte from register to memory
	load linked word	`ll $s1,20($s2)`	$s1 = Memory[$s2 + 20]	Load word as 1st half of atomic swap
	store condition. word	`sc $s1,20($s2)`	Memory[$s2+20]=$s1;$s1=0 or 1	Store word as 2nd half of atomic swap
	load upper immed.	`lui $s1,20`	$s1 = 20 * 2^{16}	Loads constant in upper 16 bits
Logical	and	`and $s1,$s2,$s3`	$s1 = $s2 & $s3	Three reg. operands; bit-by-bit AND
	or	`or $s1,$s2,$s3`	$s1 = $s2 | $s3	Three reg. operands; bit-by-bit OR
	nor	`nor $s1,$s2,$s3`	$s1 = ~ ($s2 | $s3)	Three reg. operands; bit-by-bit NOR
	and immediate	`andi $s1,$s2,20`	$s1 = $s2 & 20	Bit-by-bit AND reg with constant
	or immediate	`ori $s1,$s2,20`	$s1 = $s2 | 20	Bit-by-bit OR reg with constant
	shift left logical	`sll $s1,$s2,10`	$s1 = $s2 << 10	Shift left by constant
	shift right logical	`srl $s1,$s2,10`	$s1 = $s2 >> 10	Shift right by constant
Conditional branch	branch on equal	`beq $s1,$s2,25`	if ($s1 == $s2) go to PC + 4 + 100	Equal test; PC-relative branch
	branch on not equal	`bne $s1,$s2,25`	if ($s1!= $s2) go to PC + 4 + 100	Not equal test; PC-relative
	set on less than	`slt $s1,$s2,$s3`	if ($s2 < $s3) $s1 = 1; else $s1 = 0	Compare less than; for `beq`, `bne`
	set on less than unsigned	`sltu $s1,$s2,$s3`	if ($s2 < $s3) $s1 = 1; else $s1 = 0	Compare less than unsigned
	set less than immediate	`slti $s1,$s2,20`	if ($s2 < 20) $s1 = 1; else $s1 = 0	Compare less than constant
	set less than immediate unsigned	`sltiu $s1,$s2,20`	if ($s2 < 20) $s1 = 1; else $s1 = 0	Compare less than constant unsigned
Unconditional jump	jump	`j 2500`	go to 10000	Jump to target address
	jump register	`jr $ra`	go to $ra	For switch, procedure return
	jump and link	`jal 2500`	$ra = PC + 4; go to 10000	For procedure call

Note how just about all operations have exactly three operands. This conforms to the philosophy of keeping the hardware simple: hardware for a variable number of operands is more complicated than hardware for a fixed number.

SUBSECTION 2.3

Operands of the Computer Hardware

2.3.1 Memory Operands

Let's say we have the following C statement

```
1  g = h + A[8];
```

What will be the associated MIPS code if *g* and *h* are in registers $s1 and $s2, and that the base address of the array is in $s3.

```
1  lw    $t0, 32($s3)    # Temporary reg $t0 gets A[8]
2  add   $t0, $s2, $t0   # Temporary reg $t0 gets h + A[8]
3  sw    $t0, 48($s3)    # A[12] <—— $t0
```

Note that MIPS uses byte addressing, and so to get to the 8th index, you need to add $8 * 4$ since the size of each array index is 4 bytes. Also note that words must start with addresses that are multiples of 4.

Alignment Restriction is the requirement that data be aligned in memory on natural boundaries.

2.3.2   Constant or Immediate Operands

Let's say for example that we want to add 5 to some register for whatever reason, instead of loading that from a memory location into a temporary register, and then adding the temporary register to the desired register we can use the instruction add immediate:

```
1   addi    $s3, $s3, 4       # $s3 = $s3 + 4
```

By including this constant opperation the processor can operations much faster and using less energy. Because more than half of MIPS arithmetic instructions have a constant as an operand when running the SPEC CPU2006 benchmarks this is an example of making the common case fast.

Subsection 2.4
Signed and Unsigned Numbers

Overflow occurs when after performing an arithmatic operation on two numbers results in a number that can't be stored in the number of bits available to register (32 in case of MIPS).

2.4.1   Unsigned Numbers

In any number base, the value of the ith digit d is

$$d \times \text{ Base}^i.$$

Thus for example, in binary

$$\begin{aligned}
101 &= (1 \cdot 2^2) + (0 \cdot 2^1) + (1 \cdot 2^0) \\
&= 4 + 0 + 1 \\
&= 5.
\end{aligned}$$

2.4.2   Twos Complement

The idea of twos complement is that the most significant bit indicates the sign of the number, 1 if negative and 0 if positive (zero being treated as a positive number). But instead of merely reprsenting the sign, it represents the negative equivalent value of that number. So for example

$$\begin{aligned}
1 \text{ represents } &-1 \\
10 \text{ represents } &-2 \\
100 \text{ represents } &-4
\end{aligned}$$

and the other bits are their normal positive values:

$$\begin{aligned}
11 \text{ represents } &-1 \\
111 \text{ represents } &-1 \\
101 \text{ represents } &-3 \\
110 \text{ represents } &-2.
\end{aligned}$$

To achieve two's complement representation you

1. Start with the equivalent positive number

2. Invert all bits (change every 0 to 1, and every 1 to 0)

3. Add 1 to the inverted number, ignoring overflow.

The reason this works is because $x + \bar{x} = -1$ and therefore $\bar{x} + 1 = -x$, where $\bar{x}$ is x inverted. You can use this process to convert from negative to positive and vice versa:

- Twos complement representation of -5

$$0101 = 5$$
$$1010 = \bar{5} \qquad \textit{Inverse of 5}$$
$$1011 = -5 \qquad \textit{-5 represented in twos complement}$$

- Twos complement representation of 5

$$1011 = -5$$
$$0100 = \overline{-5}$$
$$0101 = 5.$$

In two's complement, if you want to use more bits to represent the same number, you just use **sign extension** (repeat the most significant bit):

$$0111 \equiv 7 \equiv 0000\ 0111$$
$$1011 \equiv -5 \equiv 1111\ 1011.$$

SUBSECTION 2.5
Representing Instructions in the Computer

Now that we know how to tell the computer what to do (assembly language), what instructions are the computer actuall following (machine language)? Well, first of all, computers just read a stream of 1's and 0's. Depending upon the context, this stream could represent a string, a number, or an instruction.

2.5.1   Hexadecimal

Although we now know how to write numbers in binary, it can be a bit tedious, so we use hexadecimial (base 16). Because both binary and hexadecimal are powers of two, they play nicely together.

Hexadecimal	Binary	Hexadecimal	Binary	Hexadecimal	Binary	Hexadecimal	Binary
0_{hex}	0000_{two}	4_{hex}	0100_{two}	8_{hex}	1000_{two}	c_{hex}	1100_{two}
1_{hex}	0001_{two}	5_{hex}	0101_{two}	9_{hex}	1001_{two}	d_{hex}	1101_{two}
2_{hex}	0010_{two}	6_{hex}	0110_{two}	a_{hex}	1010_{two}	e_{hex}	1110_{two}
3_{hex}	0011_{two}	7_{hex}	0111_{two}	b_{hex}	1011_{two}	f_{hex}	1111_{two}

Figure 6. To convert from hexadecimal to binary just replace hexadecimal digit by the corresponding four binary digits and vice versa. If the length of the binary number is not a multiple o 4, go from right to left.

2.5.2   MIPS Fields

Because computers just read a string of 1's and 0's, how does the computer know where each instruction begins and ends? When converting from assembly $\longrightarrow$ binary, MIPS instructions are formatted into 32 bits.
Register Instructions (R-type)

op	rs	rt	rd	shamt	funct
6 bits	5 bits	5 bits	5 bits	5 bits	6 bits

Figure 7. The various MIPS fields for R-type (register) instructions.

- op: Basic operation of the instruction, traditionally called the opcode. The value of this lets computer know meaning of, and size of the following fields.

- rs: The first register source operand

- rt: The second register source operand

- rd: The register destination operand. It gets the result of the operation.

- shamt: Shift amount. See 2.6 for more details.

- funct: Function. This field, often called the function code, selects the specific variant of the operation in the op field.

Immediate Instructions (I-type)

You might have noticed that not all instructions fit into the above format. For example Let's say you wanted to add a constant to a register and store it in another? Well, at the moment we would only have 5-bits available to hold a number, that's not a lot. The max constant value would be $2^5 - 1 = 31$ (assuming unsigned number), that's not very big. Instead we have a different instruction format:

op	rs	rt	constant or address
6 bits	5 bits	5 bits	16 bits

Note that instructions with this format type are not just add immidiate type instructions, they are also used for load word instructions (among others). The following instruction

```
1   lw    $t0, 1200($t1)   # $t0 <—— A[300]
2   add   $t0, $s2, $t0    # $t0 <—— h + A[300]
3   sw    $t0, 1200($t1)   # A[300] <—— $t0
```

is expressed as follows:

Op	rs	rt	rd	address/shamt	funct
35	9	8		1200	
0	18	8	8	0	32
43	9	8		1200	

and in binary:

100011	01001	01000	0000 0100 1011 0000		
000000	10010	01000	01000	00000	100000
101011	01001	01000	0000 0100 1011 0000		

From this, you can see that each register has an associated reference number:

NAME	NUMBER	USE	PRESERVED ACROSS A CALL?
$zero	0	The Constant Value 0	N.A.
$at	1	Assembler Temporary	No
$v0-$v1	2-3	Values for Function Results and Expression Evaluation	No
$a0-$a3	4-7	Arguments	No
$t0-$t7	8-15	Temporaries	No
$s0-$s7	16-23	Saved Temporaries	Yes
$t8-$t9	24-25	Temporaries	No
$k0-$k1	26-27	Reserved for OS Kernel	No
$gp	28	Global Pointer	Yes
$sp	29	Stack Pointer	Yes
$fp	30	Frame Pointer	Yes
$ra	31	Return Address	Yes

Here are the opcodes for the instructions we have seen so far:

instruction	Format	op	rs	rt	rd	shamt	funct	address
add	R	0	reg	reg	reg	0	32_{ten}	n.a.
sub (subtract)	R	0	reg	reg	reg	0	34_{ten}	n.a.
add immediate	I	8_{ten}	reg	reg	n.a.	n.a.	n.a.	constant
lw (load word)	I	35_{ten}	reg	reg	n.a.	n.a.	n.a.	address
sw (store word)	I	43_{ten}	reg	reg	n.a.	n.a.	n.a.	address

2.5.3   Check Yourself

What MIPS instruction does this represent?

op	rs	rt	rd	shamt	funct
0	8	9	10	0	34

From the ocode of 0 and function code of 34 we know this is a subtract function call.

We are dealing with registers $8, 9, 10$ which are `$t0`, `$t1`, and `$t2` respectively. Because the MIPS register arguments are orderd **rd**, **rs**, **rt** we have get the MIPS code

```
1   sub     $t2, $t0, $t1
```

SUBSECTION 2.7
Instructions for Making Decisions

2.7.1   Branch Instructions

MIPS has two decision making instructions:

- `beq register1, register2, L1`
  Go to the statement labeled `L1` if the value in `register1` equals the value in `register2`;

- `bne register1, register2, L1`
  Go to statement labeled `L1` if value in `register1` does not equal value in `register2`.

For example, the C code

```
1   if (i == j)
2       f = g + h;
3   else
4       f = g−h;
```

where `f` through `j` correspond to `$s0` through `$s4`, corresponds to the MIPS code

```
1       bne     $s3, $s4, Else
2       add     $s0, $s1, $s2     # f = g + h
3       j       Exit              # go to Exit
4
```

12

```
5  Else :
6      sub      $s0 , $s1 , $s2      # f = g−h
7
8  Exit :
```

Notice how we did things in reverse; instead of checking for equality (like in the C code), we checked for inequality.

The assembler relieves the compiler and the assembly language programmer from the tedium of calculated addresses for branches.

2.7.2   Loops

How could we use our branch statements to compile the following while Loop:

```
1  while ( save [ i ] == k)
2      i += 1;
```

where `i` and `k` correspond to registers `$s3` and `$s5`. The address of `save[0]` is in `$s6`.

```
1  Loop :
2    sll     $t1 , $s3 , 2      # $t1 <—— i * 4
3    add     $t1 , $t1 , $s6    # $t1 <—— address of save [ i ]
4    lw      $t0 , 0( $t1 )     # $t0 <—— save [ i ]
5    bne     $t0 , $s5 , Exit   # if ( save [ i ] != k) Exit
6    addi    $s3 , $s3 , 1      # i += 1
7    j       Loop               # go to Loop
8
9  Exit :
```

Because MIPS is byte addressed, to get the memory address of `save[i]` we need to multiply $i \cdot 4$ and add that to the base address. We then load that value to a temporary register, and test the condition.

This is an example of a **basic block**. This is a sequence without branches, except possibly at the end, and without branch targets or branch labels, except possibly at the beginning.

2.7.3   Unsigned Comparison Tests

In addition to testing for equality, we might also want to test for less than, or greater than. To test if less than the instruction "set on less than"

```
1    sltu      $t0 , $s3 , $s4    # $t0 = 1 if $s3 < $s4 ; otherwise 0
```

This sets the destination register to 1 if less than, and 0 if not. An immediate version of this instruction also exists:

```
1    sltiu     $t0 , $s2 , 10    # $t0 = 1 if $s2 < 10 , 0 otherwise
```

All relative conditions equal, not equal, less than, less than or equal, greater than, and greater than or equal, can all be created using `slt`, `slti`, `beq`, `bne`, and `$zero` register.

2.7.4   Signed Comparison Tests

We have to remember that with negative numbers, the most significant bit is turned on, and so using an unsigned comparison between -1 and 5: $1111_{\text{TWO}} < 0101_{\text{TWO}}$ would return false. Thus, for signed comparisons we have `slt` and `slti`:

```
1    slt    $t0, $s3, $s4    # $t0 = 1 if $s3 < $s4; otherwise 0
2    slti   $t0, $s2, 10     # $t0 = 1 if $s2 < 10, 0 otherwise
```

2.7.5   Bounds Checking Shortcut

Let's say we want to check if both

$$x < y \quad \text{and} \quad x \geq 0$$

where x is a and y are signed integers (and $y > 0$). Assume that `$s0` holds x and `$s1` holds y. Than we could do

```
1   # Check if x < y
2      slt    $t0, $s0, $s1
3      beq    $t0, $zero, False
4
5   # Check if x >= 0
6      slt    $t0, $s0, $zero        # If x < 0
7      bne    $t0, $zero, False
8
9   # Statement is true
10     j       True
```

We can actually check for both $x < y$ and $x \geq 0$ in one go by treating x and y as unsigned integers. The idea here is that since a negative number will have its most significant bit turned on, it looks like a very large unsigned number. Therefore no matter how large y is, if x is a negative number it will appear as a larger unsigned number. Thus,

$$0 \leq x \leq y$$

will only be true if x is nonnegative and less than y, and when treated as unsigned numbers, testing if $x < y$ will also check if $x \geq 0$:

```
1   sltu $t0, $s0, $s1
2   beq  $t0, $zero, False
3   j     True
```

This shortcut is really helpful if we want to check if some index is out of bounds for a given array. For example, let's say we have an array `example` whose upper bounds y. Then we could use the code above to check if some index i is in bounds ($0 \leq x < y$).

SUBSECTION 2.8
Supporting Procedures in Computer Hardware

Thus far, we have only really learned how to perform a sequence of instructions, but what if we want to perform that same sequence of instructions a lot of times, and maybe we want to act upon some variables. In C we have functions, and in assembly, we have **procedures**. To perform a procedure in MIPS, we follow the follwoing steps:

A **procedure** is a stored subroutine that performs a specific task based on the parameters with which it is provided.

1. Put parameters in a place where the procedure can access them.

2. Transfer control to the procedure.

3. Acquire the storage resources needed for the procedure.

4. Perform the desired task.

5. Put the result value in a place where the calling program can access it.

6. Return control to the point of origin, since a procedure can be called from several points in a program.

The following registers are used for procedure calling

- `$a0-$a3`: four argument registers in which to pass parameters

- `$v0-$v1`: two value registers in which to return values

- `$ra`: one return address register to return to the point of origin

The return address register is populated by the **jump-and-link-instruction** (jal):

`jal ProcedureAddress`. This instruction simultaneously jumps to ProcedureAddress, and populates the return address register (`$ra`) with the address of the following instruction. This way the computer knows what to perform once the procedure is finished.

The **caller** is the one that puts the parameter values in `$a0-$a3` and uses `jal X` to jump to procedure `X` (**callee**). Once the callee has finished, it places the results in `$v0` and `$v1`, and then returns control to the caller using `jr $ra`. This is an unconditional jump to the address specified in a register.

To make all of this work, we need to store the address of the current instruction being executed. We do this in a register called the **program counter** (PC). The `jal` instruction saves PC + 4 in register `$ra`.

The **return address** is a link to the calling site that allows a procedure to return to the proper address; in MIPS it is stored in regeister `$ra`.

2.8.1 Nested Procedures

Let's say you have the following program:

```c
 1  #include <stdio.h>
 2
 3  int fact(int n)
 4  {
 5      if (n < 1)
 6          return 1;
 7      else
 8          return (n * fact(n-1));
 9  }
10
11  int main() {
12      int result = fact(3);
13      printf("%d\n", result);
14      return 0;
15  }
```

Because `fact(n)` can call itself, you have to be careful about saving the return address register (`$ra`), as well as any other registers that may be used (more info below). We can write the above program in MIPS as follows:

```
1   main:
2      addi      $a0, $zero, 3
3      jal       fact
4      add       $a0, $v0, $zero # store result for printing
5      addi      $v0, $zero, 1   # print int syscall
6      syscall
7      addi      $v0, $zero, $10 # exit syscall
8      syscall
9
10  fact:
11     addi      $sp, $sp, -8     # make space for two items on stack
12     sw        $ra, 4($sp)      # save return address
13     sw        $a0, 0($sp)      # save arg n
14     slti      $t0, $a0, 1      # test if n < 1
15     beq       $t0, $zero, L1   # if n >= 1, goto L1
16     lw        $ra, 4($sp)      # restore $ra
17     lw        $a0, 0($sp)      # restore arg n
18     addi      $v0, $zero, 1    # return 1
19     addi      $sp, $sp, 8      # pop two items from stack
20     jr        $ra              # return to caller
21
22  L1:
23     addi      $a0, $a0, -1      # n += -1
24     jal       fact              # calls fact and $ra <-- PC+4
25     lw        $a0, 0($sp)       # restore arg n
26     lw        $ra, 4($sp)       # restore $ra
27     addi      $sp, $sp, 8       # Pops two items from stack
28     mul       $v0, $a0, $v0     # return n*fact(n-1)
29     jr        $ra               # return to caller
```

Let's look at how the register values, and stack changes as we go through the program:

1. In `main` (before first call of fact):

   ```
   $a0 <-- 3
   $ra <-- return to main
   ```

2. During first call of `fact(3)`

```
                Stack
          _____
         |    ra  of  main   |
          _____   <-- -4
         |         3         |
          _____   <-- -8
```

16

3. First break to `L1`:

```
$a0  <—— 3−1 = 2                    Stack
$ra  <—— return  to  L1      ————————————————
                             |   ra  of  main  |
                             ————————————————   <—— −4
                             |        3         |
                             ————————————————   <—— −8
```

4. During second call to `fact(2)`

```
$a0  =  3−1 = 2                     Stack
$ra  =  return  to  L1       ————————————————
                             |   ra  of  main  |
                             ————————————————   <—— −4
                             |        3         |
                             ————————————————   <—— −8
                             |   ra  of  L1    |
                             ————————————————   <—— −12
                             |        2         |
                             ————————————————   <—— −16
```

5. During second break to `L1`

```
$a0  <—— 2−1 = 1                    Stack
$ra  <—— return  to  L1      ————————————————
                             |   ra  of  main  |
                             ————————————————   <—— −4
                             |        3         |
                             ————————————————   <—— −8
                             |   ra  of  L1    |
                             ————————————————   <—— −12
                             |        2         |
                             ————————————————   <—— −16
```

6. During third call to `fact(1)`

```
$a0  =  1                    Stack
$ra  =  return  to  L1       ————————————————
                             |   ra  of  main  |
                             ————————————————   <—— −4
                             |        3         |
                             ————————————————   <—— −8
                             |   ra  of  L1    |
                             ————————————————   <—— −12
                             |        2         |
                             ————————————————   <—— −16
                             |   ra  of  L1    |
                             ————————————————   <—— −20
                             |        1         |
                             ————————————————   <—— −24
```

7. During third break to `L1`

```
$a0 <——  1−1 = 0                   Stack
$ra <——  L1
                            ———————————————
                           |   ra  of  main   |
                            ———————————————  <—— −4
                           |        3         |
                            ———————————————  <—— −8
                           |    ra  of  L1    |
                            ———————————————  <—— −12
                           |        2         |
                            ———————————————  <—— −16
                           |    ra  of  L1    |
                            ———————————————  <—— −20
                           |        1         |
                            ———————————————  <—— −24
```

8. During fourth call to `fact(0)`

```
$a0 = 0                              Stack
$ra = L1
                            ———————————————
                           |   ra  of  main   |
$ra <—— stack[−28] = L1     ———————————————  <—— −4
$a0 <—— stack[−32] = 0     |        3         |
$v0 <—— 1 * 1 = 1           ———————————————  <—— −8
                           |    ra  of  L1    |
                            ———————————————  <—— −12
                           |        2         |
                            ———————————————  <—— −16
                           |    ra  of  L1    |
                            ———————————————  <—— −20
                           |        1         |
                            ———————————————  <—— −24
                           |    ra  of  L1    |
                            ———————————————  <—— −28
                           |        0         |
                            ———————————————  <—— −32$
```

9. Backtracking to 3rd `L1`:

$ra <—— stack$[-20]$ = L1   Stack
$a0 <—— stack$[-24]$ = 1   ——————————
$v0 <—— 1 * 1 = 1 | ra of main |
                                                         <—— -4
                                  | 3 |
                                                         <—— -8
                                  | ra of L1 |
                                                         <—— -12
                                  | 2 |
                                                         <—— -16
                                  | ra of L1 |
                                                         <—— -20
                                  | 1 |
                                                         <—— $-24$$

10. Backtracking to 2nd `L1`:

$ra <—— stack$[-12]$ = L1   Stack
$a0 <—— stack$[-16]$ = 2   ——————————
$v0 <—— 1 * 2 = 2 | ra of main |
                                                         <—— -4
                                  | 3 |
                                                         <—— -8
                                  | ra of L1 |
                                                         <—— -12
                                  | 2 |
                                                         <—— $-16$$

11. Backtracking to 1st `L1`:

$ra <—— stack$[-4]$ = main   Stack
$a0 <—— stack$[-8]$ = 3   ——————————
$v0 <—— 2 * 3 = 6 | ra of main |
                                                         <—— -4
                                  | 3 |
                                                         <—— $-8$$

12. Back in main

$a0 <—— 6
$v0 <—— 1
Print 1
Exit

If you examine things closely, you can actually forgoe the restoring of `$ra` and `$a0` inside of fact, since those remain unchanged.

2.8.2   Calling Conventions

Rules for all calls

$sp→ 7fff fffc$_{hex}$ Stack
                                            ↓

                                            ↑
                                     Dynamic data
$gp→ 1000 8000$_{hex}$ Static data
         1000 0000$_{hex}$
                                     Text
pc→ 0040 0000$_{hex}$
                                     Reserved
                 0

Figure 8. How the memory is split up.

- The first instruction in the function manipulates `$sp` to allocate F words of stack:

$$\texttt{addi \$sp, \$sp, -F*4}$$

where $F = A + L + P + (1 \text{ for } \texttt{\$ra} \text{ if needed }) + (1 \text{ for } \texttt{\$fp} \text{ if in use }) + (2 \text{ for } \texttt{\$gp}) + $ padding . Also, F is the frame size in words, and should be even.
Corollary 1: `$sp` is not changed at any other point than the first and last instructions in the function.

- If you want to use any preserved register in any place during the function, you must store them into the stack when you **enter** the function (not when you use it). Let's assume we are saving `$s0`, `$s1`, and `$s2` registers, as well as `$ra` and `$sp`. Then we would write:

```
sw  $ra ,  (F−1)∗4($sp)     # Not needed for leaf functions
sw  $fp ,  (F−2)∗4($sp)
sw  $s2 ,  (F−3)∗4($sp)
sw  $s1 ,  (F−4)∗4($sp)
sw  $s0 ,  (F−1)∗4($sp)
```

- After saving preserved registers, copy `$sp` into `$fp`, so that from now until you start returning, you can use offsets from `$fp` to access the locations. Save and restore local variables as needed after this point, using `$fp` and the spots already allocated for them on the stack.

- You may pass information **into** a function only through `$a0-$a3` and extra stack arguments. As the called function, you may not use the value held in any other register except `$sp`, `$fp`, `$gp` and `$ra` (i.e. don't try and use `$v0` or `$s0`). The 4 (fifth) argument to the current function is available at $(F+4) * 4(\texttt{\$fp})$. You may store argument 0 from `$a0` into its reserved spot at $F * 4(\$fp)$, arg1 into $(F+1) * 4(\texttt{\$fp})$, etc., if you want.

- You may pass information **out of** a function through only `$v0-$v1`.

- You may only enter a function at its beginning. You may not `jal` into the middle of a function.

- You may have only one `jr $ra` in each function. It is the last instruction in the function, following the instructions that restore the registers saved.

Extra rules for non-leaf functions:

- Look at the declarations for all procedures this function *might* call. Find the one with the largest number of arguments, A, and then make A at least 4. You must allocate $A \cdot 4$ bytes on the stack upon entry into your function.

- If passing more than 4 arguments, then place 4 arguments into `$a0-$a3`. Every argument past the 4^{th} argument is placed on the stack. For example suppose 6 arguments are in `$s0-$s5`:

```
move $a0 ,  $s0     # 0($fp) reserved on the stack
move $a1 ,  $s1     # 4($fp) reserved on the stack
move $a2 ,  $s2     # 8($fp) reserved on the stack
move $a3 ,  $s3     # 12($fp) reserved on the stack
sw  $s4 ,  16($fp)  # 5th arg in the stack. Always here!
sw  $s5 ,  20($fp)  # 5th arg in the stack. Always here!
```


Figure 9. Stack organization during call

- You must preserve register `$ra` because a called-function `jal` destroys the value, and it is a preserved register.

SUBSECTION 2.10
MIPS Addressing for 32-bit Immediates and Addresses

2.10.1   32-bit Constant into Register

If you have been following closely, you may have noticed that when adding a constant to a register there are only 16 bits available for the constant (as it is an I-type instruction):

```
 _____ _____ _____ _____
|  op   |  rt  |  rs  |      constant    |
 _____ _____ _____ _____
    6       5      5            16
```

The question then remains, what do you do if you want to place a constant larger than 16 bits? You achieve this with the following two instructions:

1. `lui rt, immediate`
   which loads immediate value into the upper 16 bits of `rt` and zeros lower 16 bits.

2. `ori rt, rs, immediate`
   Bitwise or of `rs` and `immediate` placed into `rt` (`rt <- rs | imm`).

A word of caution. Using the instruction `addi` sign extends the msb of immediate.

2.10.2   Branch Addressing

Branch instructions take are I-type instructions:

```
 _____ _____ _____ _____
|  op   |  rs  |  rt  |encoded  address  |
 _____ _____ _____ _____
    6       5      5            16
```

The target of the encoded addres is

$$\text{Target} = 4(\text{encoded address}) + \texttt{PC} + 4.$$

That is, branch instructions are word addressed (4(encoded address)), and relative to the next instruction ($\texttt{PC} + 4$).

2.10.3   Jump Addressing

Jump instructions are J-type instructions:

```
 _____ _____
|  op   |        encoded  address       |
 _____ _____
    6                  26
```

In this case the lower 28 bits of the address to be jumped to are

$$4 \cdot \text{encoded address}$$

and the upper four are that of the `PC`. If you need to jump to a really far away address (more than 64 million instructions away), then you use jump register instruction (`jr`).