

CS10 – COMPUTER ARCHITECTURE AND ORGANIZATION

MY NOTES

LEONARD MOHR

Student at Foothill College

leonardmohr@gmail.com

Contents

1 Computer Abstractions and Technology	1
1.1 Introduction	1
1.1.1 Program Performance	1
1.2 Eight Great Ideas in Computer Architecture	2
1.3 Below Your Program	3
1.4 Under the Covers	3
1.4.1 Parts of the Computer	3
1.6 Performance	4
1.6.1 Clock Rate and Clock Period	4
1.6.2 Instruction Performance	5
1.6.3 The Classic CPU Performance Equation	5
1.10 Fallacies and Pitfalls	5
1.10.1 Amdahl's Law	5
1.10.2 MIPS	6
1.10.3 Check Yourself	6
2 Instructions: Language of the Computer	7
2.2 Operations of the Computer Hardware	7
2.2.1 MIPS Operands	7
2.2.2 MIPS Instructions	8
2.3 Operands of the Computer Hardware	8
2.3.1 Memory Operands	8
2.3.2 Constant or Immediate Operands	9
2.4 Signed and Unsigned Numbers	9
2.4.1 Unsigned Numbers	9
2.4.2 Twos Complement	9
2.5 Representing Instructions in the Computer	10
2.5.1 Hexadecimal	10
2.5.2 MIPS Fields	10
2.5.3 Check Yourself	12
2.7 Instructions for Making Decisions	12
2.7.1 Branch Instructions	12
2.7.2 Loops	13
2.7.3 Unsigned Comparison Tests	13
2.7.4 Signed Comparison Tests	13
2.7.5 Bounds Checking Shortcut	14
2.8 Supporting Procedures in Computer Hardware	14
2.8.1 Nested Procedures	15
2.8.2 Calling Conventions	19
2.10 MIPS Addressing for 32-bit Immediates and Addresses	21
2.10.1 32-bit Constant into Register	21

2.10.2	Branch Addressing	21
2.10.3	Jump Addressing	21
2.12	Translating and Starting a Program	22
2.12.1	Compiler	22
2.12.2	Assembler	22
2.12.3	Linker	22
2.12.4	Loader	23
2.14	Arrays versus Pointers	23
2.14.1	Array	23
3	Arithmetic for Computers	24
3.2	Addition and Subtraction	24
3.2.1	Overflow Twos Complement	25
3.2.2	Unsigned Overflow	25
3.3	Multiplication	26
3.3.1	Faster Multiplication	27
3.3.2	Even Faster Multiplication	28
3.4	Division	28
3.4.1	Faster Division	29
4	Basics of Logic Design	29
4.1	Truth Tables	29
4.2	Boolean Algebra	30
4.2.1	Laws	30
4.3	Logic Blocks	30
4.3.1	Decoder	31
4.3.2	Multiplexer	31
4.3.3	Two-Level Logic	32
4.3.4	Programmable Logic Array	33
4.3.5	Don't Cares	33
4.3.6	Arrays of Logic Units	34
5	Constructing a Basic Arithmetic Logic Unit	35
5.1	1-Bit ALU	35
5.1.1	AND / OR Logical Operation	35
5.1.2	Addition	35
5.1.3	Completing the 1-Bit ALU	36
5.2	32-Bit ALU	36
5.2.1	Subtraction	36
5.2.2	NOR (Neither a nor b)	37
5.2.3	Shift Less Than	38
6	Memory Elements	38
6.1	S-R Latch	39
6.2	Latches	39
6.2.1	Transparent Latch	39
6.2.2	Gated Latch	39

6.2.3	Flip-Flop	40
7	The Processor	42
7.1	Introduction	42

SECTION 1

Computer Abstractions and Technology

SUBSECTION 1.1

Introduction

Decimal term	Abbreviation	Value	Binary term	Abbreviation	Value	% Larger
kilobyte	KB	10^3	kibibyte	KiB	2^{10}	2%
megabyte	MB	10^6	mebibyte	MiB	2^{20}	5%
gigabyte	GB	10^9	gibibyte	GiB	2^{30}	7%
terabyte	TB	10^{12}	tebibyte	TiB	2^{40}	10%
petabyte	PB	10^{15}	pebibyte	PiB	2^{50}	13%
exabyte	EB	10^{18}	exbibyte	EiB	2^{60}	15%
zettabyte	ZB	10^{21}	zebibyte	ZiB	2^{70}	18%
yottabyte	YB	10^{24}	yobibyte	YiB	2^{80}	21%

Figure 1. We can describe storage in binary or decimal notation. We are much more familiar with the decimal term. Note also the size difference between the two: binary term gets progressively larger.

1.1.1 Program Performance

One of the main goals for both the hardware designer and the software designer is to improve performance. This can be achieved in different ways (just think about how quick sort is much faster than bubble sort, and the apple M1 processor is much faster than the intel processors they replaced).

Hardware or software component	How this component affects performance	Where is this topic covered?
Algorithm	Determines both the number of source-level statements and the number of I/O operations executed	Other books!
Programming language, compiler, and architecture	Determines the number of computer instructions for each source-level statement	Chapters 2 and 3
Processor and memory system	Determines how fast instructions can be executed	Chapters 4, 5, and 6
I/O system (hardware and operating system)	Determines how fast I/O operations may be executed	Chapters 4, 5, and 6

Figure 2. Here we can see the various ways program performance can be improved.

Check Yourself

- As mentioned earlier, both the software and hardware affect the performance of a program. Can you think of examples where each of the following is the right place to look for a performance bottleneck?

- The algorithm chosen

If a program that sorts a really long list of names is taking a long time, you might want to look at the algorithm being used.

- The programming language or compiler

If your program could be language dependent, you could look at using a compiled language (like C), as opposed to an interpreted language like Python.

- c) The operating system

If one program runs well, but two at a time don't, then maybe the operating system isn't distributing its resources efficiently.

- d) The processor

If the computer in general is using a lot of energy, you might want to look at the processor. Similarly, if you are doing compute heavy tasks such as video editing, you might need a processor upgrade.

- e) The I/O system and devices

If it takes a long time to write to a hard drive, you might look at upgrading to a SSD.

SUBSECTION 1.2

Eight Great Ideas in Computer Architecture

1. Moore's Law

Moore's Law resulted from a 1965 prediction that integrated circuit (IC) resources would double every 18-24 months.

2. Use Abstraction to Simplify Design

To increase productivity, by design both computer architects and programmers try to use abstractions to hide the lower-level details and provide a simpler interface to work with. For example, the operating system abstracts away the complexity of the memory system so that programs are provided with a much simplified view of the memory, and the details are handled by the operating system.

3. Make the Common Case Fast

Better performance gains can be made if you optimize for what the program is going to do most often.

4. Performance via Parallelism

We can make a program a lot faster by performing multiple tasks at once. This is especially true with the advent of multi-core processors.

5. Performance via Pipelining

If you are moving a lot of bricks from one place to another (using just manpower) it would be a lot more efficient to set up a line of people, and pass the bricks down the line, than to just have everyone running back and forth. The same principle can be used in computers.

6. Performance via Prediction

When a processor encounters an if statement, it might say that on average the result is true, and proceed like it is, so that it keeps going, and then the if qualifier can be processed at a later date.

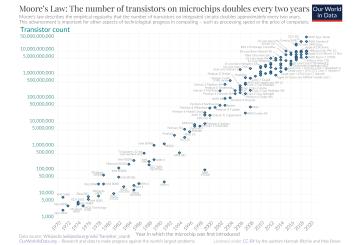


Figure 3. Moore's Law in action.

7. Hierarchy of Memories

By using different types of memory, from really small and really fast, to really large and slow, a memory hierarchy is created. Caches give the programmer the illusion that main memory is nearly as fast as the top of the hierarchy and nearly as big and cheap as the bottom of the hierarchy.

8. Dependability via Redundancy

Because we are sad when a computer dies, we want to prevent the death of the computer. To do this we make them dependable. This can be achieved by including redundant components that can both take over in the event of a failure as well as detect when a failure has occurred.

SUBSECTION 1.3

Below Your Program

The Operating System is a great example of abstraction. Some of its most important functions are:

- Handling basic input and output operations
- Allocating storage and memory
- Provided for protected sharing of the computer among multiple applications using it simultaneously.

Another example of abstraction is high-level programming languages like C. When computers first came about, programmers programmed in binary; they wrote their programs in 1's and 0's (the language of the computer). Since that was tedious, they invented the assembler which would convert assembly language into binary code. Next came the compiler which would convert higher order languages into assembly.

Another benefit of the programming languages is it allows one to chose the specific language that is best for the task. Also, programs don't have to be written for a specific processor, since the compiler and assembler can package it for different computers.

SUBSECTION 1.4

Under the Covers

The five classic components of a computer are input, output, memory, datapath, and control, with the last two sometimes combined and called the processor. This organization is independent of hardware technology: you can place every piece of every computer, past and present, into one of these five categories.

1.4.1 Parts of the Computer

- Integrated Circuit: Also called a chip. A device combining dozens to millions of transistors.
- Central Processing Unit (CPU): Also called the processor. The active part of the computer, which contains the datapath and control and which adds numbers, tests numbers, signals I/O devices to activate, and so on.
- Datapath: The component of the processor that performs arithmetic operations
- Control: The component of the processor that commands the datapath, memory, and I/O devices according to the instructions of the program.
- Memory: The storage area in which programs are kept when they are running and that contains the data needed by the running programs.

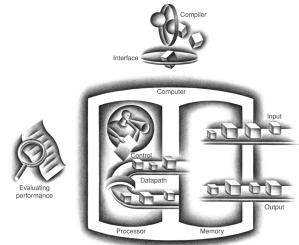


Figure 4. Here we see the flow of information in a computer. The processor gets instructions and data from memory. Input writes data to memory, and output reads data from memory. Control sends the signals that determine the operations of the datapath, memory, input, and output.

- Dynamic Random Access Memory (DRAM): Memory built as an integrated circuit; it provides random access to any location. Access times are 50 nanoseconds and cost per gigabyte in 2012 was \$5 to \$10.
- Cache Memory: Consists of a small, fast memory that acts as a buffer for the DRAM memory.
- Static Random Access Memory: SRAM is faster but less dense, and hence more expensive, than DRAM.
- Instruction Set Architector: The interface between the hardware and the lowest-level software. This includes all the information necessary to write a machine language program that will run correctly, including instructions, registers, memory access, I/O, and so on.
- Application Binary Interface (ABI): The user portion of the instruction set plus the operating system interfaces used by application programmers. It defines a standard for binary portability across computers.

SUBSECTION 1.6

Performance

When talking about computers, we often want to look at the performance of the computer. But how can we define performance?

Definition 1**Response Time**

Also referred to as **execution time**. This is the total time required for the computer to complete a task, including disk access, memory accesses, I/O activities, operating system overhead, CPU execution time, and so on.

Definition 2**Throughput / Bandwidth**

This measures the number of tasks computed per unit time.

For the time being, we will mostly be looking at response time. So that a higher number represents a machine with better performance we will say that

$$\text{Performance}_X = \frac{1}{\text{Execution time}_X}.$$

Also, we often want to say that computer “X is n times as fast as Y”, to compute n we say:

$$\frac{\text{Performance}_X}{\text{Performance}_Y} = \frac{\text{Execution Time}_Y}{\text{Execution Time}_X} = n.$$

CPU execution time or simply **CPU time** is the actual time the CPU spends computing for a specific task.

1.6.1 Clock Rate and Clock Period

Designers refer to the length of a **clock period** both as the time for a complete clock cycle (e.g., 250 picoseconds) and as the **clock rate** (e.g., 4 gigahertz, or 4GHz), which is the inverse of the clock period.

For example a clock rate of

$$\begin{aligned} 4\text{GHz} &= 4,000,000,000 \frac{\text{cycles}}{\text{sec}} \\ &\equiv \frac{1 \text{ cycle}}{4,000,000,000 \text{ second}} \\ &= 0.0000000025 \frac{\text{cycle}}{\text{second}} \\ &= \frac{1}{250} \frac{\text{cycle}}{\text{picosecond}}. \end{aligned}$$

We can relate clock cycles and clock cycle time to CPU time:

$$\text{CPU execution time for a program} = \frac{\text{CPU clock cycles for a program}}{\text{Clock cycle time}}$$

or alternatively,

$$\text{CPU execution time for a program} = \frac{\text{CPU clock cycles for a program}}{\text{Clock rate.}}$$

1.6.2 Instruction Performance

In addition to thinking about clock cycles and CPU time, we also need to think about number of instructions there are in a program and the time each instruction takes:

$$\text{CPU clock cycles} = \text{Instructions for a Program} \times \frac{\text{average clock cycles}}{\text{instruction}}$$

Clock Cycles per Instruction (CPI) is the average number of clock cycles per instruction for a program or program fragment.

1.6.3 The Classic CPU Performance Equation

Putting everything from above together we have

$$\text{CPU time} = \text{Instruction count} \times \text{CPI} \times \text{Clock cycle time}$$

and

$$\text{CPU time} = \frac{\text{Instruction count} \times \text{CPI}}{\text{Clock rate}}.$$

We also have

$$\text{Time} = \frac{\text{Seconds / Program}}{\text{Program}} = \frac{\text{Instructions}}{\text{Program}} \times \frac{\text{Clock cycles}}{\text{Instruction}} \times \frac{\text{Seconds}}{\text{Clock cycle.}}$$

SUBSECTION 1.10

Fallacies and Pitfalls

1.10.1 Amdahl's Law

Definition 3

Amdahl's Law

A rule stating that the performance enhancement possible with a given improvement

Processors these days can vary their clock rates. For example Intel Core i7 chips temporarily increase clock rate by 10% until the chip gets too warm. Thus we need to use the average clock rate for a program.

Figure 5.

is limited by the amount that the improved feature is used:

$$\text{Execution time after improvement} = \frac{\text{Execution time affected by improvement}}{\text{Amount of improvement}} + \text{Execution time unaffected}$$

where Amount of improvement = n .

We thus see that there is only so much benefit that can be achieved by improving a part of the program that is rarely used. For this reason we want to make the common case fast!

1.10.2 MIPS

One should be wary about using only a subset of the performance equation as a performance metric; one can't determine performance just by looking at clock rate, instruction count, or CPI alone.

An alternative to time is **MIPS (million instructions per second)**:

$$\text{MIPS} = \frac{\text{Instruction count}}{\text{Execution time} \times 10^6}.$$

There are a couple problems with MIPS:

- MIPS specifies the instruction execution rate but does not take into account the capabilities of these instructions. Therefore we can't compare computers with different instruction sets.
- MIPS varies between programs on the same computer; thus a computer cannot have a single MIPS rating.

$$\text{MIPS} = \frac{\text{Instruction count}}{\frac{\text{Instruction count} \times \text{CPI}}{\text{Clock rate}} \times 10^6} = \frac{\text{Clock rate}}{\text{CPI} \times 10^6}.$$

1.10.3 Check Yourself

Consider the following performance measurements for a program:

Measurement	Computer A	Computer B
Instruction count	10 billion	8 billion
Clock rate	4 GHz	4 GHz
CPI	1.0	1.1

Which computer is faster and which has the higher MIPS rating?

First looking at computer A.

$$\begin{aligned} \text{Time}(A) &= \frac{\text{CPU Clock Cycles}}{\text{Clock Rate}} \\ &= \frac{\text{Instruction Count} \times \text{CPI}}{\text{Clock Rate}} \\ &= \frac{10 \times 10^9 \text{ instructions} \times \frac{1 \text{ cycle}}{\text{instruction}}}{4 \times 10^9 \frac{\text{cycles}}{\text{second}}} \\ &= 2.5 \text{ seconds.} \end{aligned}$$

$$\begin{aligned}
 \text{MIPS}(A) &= \frac{\text{Instruction Count}}{\text{Execution Time} \times 10^6} \\
 &= \frac{10 \times 10^9 \text{ instructions}}{2.5 \text{ seconds} \times 10^6} \\
 &= \frac{4 \times 10^3 \text{ million instructions}}{\text{second}}.
 \end{aligned}$$

Now looking at computer B.

$$\begin{aligned}
 \text{Time}(B) &= \frac{\text{CPU Clock Cycles}}{\text{Clock Rate}} \\
 &= \frac{\text{Instruction Count} \times \text{CPI}}{\text{Clock Rate}} \\
 &= \frac{8 \times 10^9 \text{ instructions} \times \frac{1.1 \text{ cycles}}{\text{instruction}}}{4 \times 10^9 \frac{\text{cycles}}{\text{second}}} \\
 &= 2.2 \text{ seconds}.
 \end{aligned}$$

$$\begin{aligned}
 \text{MIPS}(B) &= \frac{\text{Instruction Count}}{\text{Execution Time} \times 10^6} \\
 &= \frac{8 \times 10^9 \text{ instructions}}{2.2 \text{ seconds} \times 10^6} \\
 &\approx \frac{3.6 \times 10^3 \text{ million instructions}}{\text{second}}.
 \end{aligned}$$

We thus see that Computer A has a higher MIPS score but runs slower.

SECTION 2

Instructions: Language of the Computer

SUBSECTION 2.2

Operations of the Computer Hardware

2.2.1 MIPS Operands

Name	Example	Comments
32 registers	\$s0-\$s7, \$t0-\$t9, \$zero, \$a0-\$a3, \$v0-\$v1, \$gp, \$fp, \$sp, \$ra, \$at	Fast locations for data. In MIPS, data must be in registers to perform arithmetic, register \$zero always equals 0, and register \$at is reserved by the assembler to handle large constants.
2^{30} memory words	Memory[0], Memory[4], ..., Memory[4294967292]	Accessed only by data transfer instructions. MIPS uses byte addresses, so sequential word addresses differ by 4. Memory holds data structures, arrays, and spilled registers.

2.2.2 MIPS Instructions

Category	Instruction	Example	Meaning	Comments
Arithmetic	add	add \$s1,\$s2,\$s3	\$s1 = \$s2 + \$s3	Three register operands
	subtract	sub \$s1,\$s2,\$s3	\$s1 = \$s2 - \$s3	Three register operands
	add immediate	addi \$s1,\$s2,20	\$s1 = \$s2 + 20	Used to add constants
Data transfer	load word	lw \$s1,20(\$s2)	\$s1 = Memory[\$s2 + 20]	Word from memory to register
	store word	sw \$s1,20(\$s2)	Memory[\$s2 + 20] = \$s1	Word from register to memory
	load half	lh \$s1,20(\$s2)	\$s1 = Memory[\$s2 + 20]	Halfword memory to register
	load half unsigned	lhu \$s1,20(\$s2)	\$s1 = Memory[\$s2 + 20]	Halfword memory to register
	store half	sh \$s1,20(\$s2)	Memory[\$s2 + 20] = \$s1	Halfword register to memory
	load byte	lb \$s1,20(\$s2)	\$s1 = Memory[\$s2 + 20]	Byte from memory to register
	load byte unsigned	lbu \$s1,20(\$s2)	\$s1 = Memory[\$s2 + 20]	Byte from memory to register
	store byte	sb \$s1,20(\$s2)	Memory[\$s2 + 20] = \$s1	Byte from register to memory
	load linked word	ll \$s1,20(\$s2)	\$s1 = Memory[\$s2 + 20]	Load word as 1st half of atomic swap
Logical	store condition. word	sc \$s1,20(\$s2)	Memory[\$s2+20] = \$s1; \$s1=0 or 1	Store word as 2nd half of atomic swap
	load upper immed.	lui \$s1,20	\$s1 = 20 * 2 ¹⁶	Loads constant in upper 16 bits
	and	and \$s1,\$s2,\$s3	\$s1 = \$s2 & \$s3	Three reg. operands; bit-by-bit AND
	or	or \$s1,\$s2,\$s3	\$s1 = \$s2 \$s3	Three reg. operands; bit-by-bit OR
	nor	nor \$s1,\$s2,\$s3	\$s1 = ~(\$s2 \$s3)	Three reg. operands; bit-by-bit NOR
	and immediate	andi \$s1,\$s2,20	\$s1 = \$s2 & 20	Bit-by-bit AND reg with constant
	or immediate	ori \$s1,\$s2,20	\$s1 = \$s2 20	Bit-by-bit OR reg with constant
	shift left logical	sll \$s1,\$s2,10	\$s1 = \$s2 << 10	Shift left by constant
	shift right logical	srl \$s1,\$s2,10	\$s1 = \$s2 >> 10	Shift right by constant
Conditional branch	branch on equal	beq \$s1,\$s2,25	if(\$s1 == \$s2) go to PC + 4 + 100	Equal test; PC-relative branch
	branch on not equal	bne \$s1,\$s2,25	if(\$s1!= \$s2) go to PC + 4 + 100	Not equal test; PC-relative
	set on less than	slt \$s1,\$s2,\$s3	if(\$s2 < \$s3) \$s1 = 1; else \$s1 = 0	Compare less than; for beq, bne
	set on less than unsigned	sltu \$s1,\$s2,\$s3	if(\$s2 < \$s3) \$s1 = 1; else \$s1 = 0	Compare less than unsigned
	set less than immediate	slti \$s1,\$s2,20	if(\$s2 < 20) \$s1 = 1; else \$s1 = 0	Compare less than constant
	set less than immediate unsigned	sltiu \$s1,\$s2,20	if(\$s2 < 20) \$s1 = 1; else \$s1 = 0	Compare less than constant unsigned
Unconditional jump	jump	j 2500	go to 10000	Jump to target address
	jump register	jr \$ra	go to \$ra	For switch, procedure return
	jump and link	jal 2500	\$ra = PC + 4; go to 10000	For procedure call

Note how just about all operations have exactly three operands. This conforms to the philosophy of keeping the hardware simple: hardware for a variable number of operands is more complicated than hardware for a fixed number.

SUBSECTION 2.3

Operands of the Computer Hardware

2.3.1 Memory Operands

Let's say we have the following C statement

```
1 g = h + A[8];
```

What will be the associated MIPS code if *g* and *h* are in registers *\$s1* and *\$s2*, and that the base address of the array is in *\$s3*.

```
1 lw    $t0 , 32($s3)      # Temporary reg $t0 gets A[8]
2 add   $t0 , $s2 , $t0    # Temporary reg $t0 gets h + A[8]
3 sw    $t0 , 48($s3)      # A[12] <— $t0
```

Note that MIPS uses byte addressing, and so to get to the 8th index, you need to add $8 * 4$ since the size of each array index is 4 bytes. Also note that words must start with addresses that are multiples of 4.

2.3.2 Constant or Immediate Operands

Let's say for example that we want to add 5 to some register for whatever reason, instead of loading that from a memory location into a temporary register, and then adding the temporary register to the desired register we can use the instruction add immediate:

```
1 addi    $s3 , $s3 , 4      # $s3 = $s3 + 4
```

By including this constant operation the processor can operations much faster and using less energy. Because more than half of MIPS arithmetic instructions have a constant as an operand when running the SPEC CPU2006 benchmarks this is an example of making the common case fast.

SUBSECTION 2.4

Signed and Unsigned Numbers

2.4.1 Unsigned Numbers

In any number base, the value of the i th digit d is

$$d \times \text{Base}^i.$$

Thus for example, in binary

$$\begin{aligned} 101 &= (1 \cdot 2^2) + (0 \cdot 2^1) + (1 \cdot 2^0) \\ &= 4 + 0 + 1 \\ &= 5. \end{aligned}$$

Overflow occurs when after performing an arithmetic operation on two numbers results in a number that can't be stored in the number of bits available to register (32 in case of MIPS).

2.4.2 Twos Complement

The idea of twos complement is that the most significant bit indicates the sign of the number, 1 if negative and 0 if positive (zero being treated as a positive number). But instead of merely representing the sign, it represents the negative equivalent value of that number. So for example

$$\begin{aligned} 1 &\text{ represents } -1 \\ 10 &\text{ represents } -2 \\ 100 &\text{ represents } -4 \end{aligned}$$

and the other bits are their normal positive values:

$$\begin{aligned} 11 &\text{ represents } -1 \\ 111 &\text{ represents } -1 \\ 101 &\text{ represents } -3 \\ 110 &\text{ represents } -2. \end{aligned}$$

To achieve two's complement representation you

1. Start with the equivalent positive number
2. Invert all bits (change every 0 to 1, and every 1 to 0)

- 3. Add 1 to the inverted number, ignoring overflow.

The reason this works is because $x + \bar{x} = -1$ and therefore $\bar{x} + 1 = -x$, where \bar{x} is x inverted. You can use this process to convert from negative to positive and vice versa:

- Twos complement representation of -5

$$\begin{array}{ll} 0101 = 5 & \\ 1010 = \bar{5} & \text{Inverse of } 5 \\ 1011 = -5 & \text{-5 represented in twos complement} \end{array}$$

- Twos complement representation of 5

$$\begin{array}{l} 1011 = -5 \\ 0100 = \overline{-5} \\ 0101 = 5. \end{array}$$

In two's complement, if you want to use more bits to represent the same number, you just use **sign extension** (repeat the most significant bit):

$$\begin{array}{l} 0111 \equiv 7 \equiv 0000\ 0111 \\ 1011 \equiv -5 \equiv 1111\ 1011. \end{array}$$

SUBSECTION 2.5

Representing Instructions in the Computer

Now that we know how to tell the computer what to do (assembly language), what instructions are the computer actually following (machine language)? Well, first of all, computers just read a stream of 1's and 0's. Depending upon the context, this stream could represent a string, a number, or an instruction.

2.5.1 Hexadecimal

Although we now know how to write numbers in binary, it can be a bit tedious, so we use hexadecimal (base 16). Because both binary and hexadecimal are powers of two, they play nicely together.

Hexadecimal	Binary	Hexadecimal	Binary	Hexadecimal	Binary	Hexadecimal	Binary
0 _{hex}	0000 _{two}	4 _{hex}	0100 _{two}	8 _{hex}	1000 _{two}	C _{hex}	1100 _{two}
1 _{hex}	0001 _{two}	5 _{hex}	0101 _{two}	9 _{hex}	1001 _{two}	D _{hex}	1101 _{two}
2 _{hex}	0010 _{two}	6 _{hex}	0110 _{two}	A _{hex}	1010 _{two}	E _{hex}	1110 _{two}
3 _{hex}	0011 _{two}	7 _{hex}	0111 _{two}	B _{hex}	1011 _{two}	F _{hex}	1111 _{two}

Figure 6. To convert from hexadecimal to binary just replace hexadecimal digit by the corresponding four binary digits and vice versa. If the length of the binary number is not a multiple of 4, go from right to left.

2.5.2 MIPS Fields

Because computers just read a string of 1's and 0's, how does the computer know where each instruction begins and ends? When converting from assembly → binary, MIPS instructions are formatted into 32 bits.

Register Instructions (R-type)



Figure 7. The various MIPS fields for R-type (register) instructions.

- op: Basic operation of the instruction, traditionally called the opcode. The value of this lets computer know meaning of, and size of the following fields.
- rs: The first register source operand
- rt: The second register source operand
- rd: The register destination operand. It gets the result of the operation.
- shamt: Shift amount. See 2.6 for more details.
- funct: Function. This field, often called the function code, selects the specific variant of the operation in the op field.

Immediate Instructions (I-type)

You might have noticed that not all instructions fit into the above format. For example Let's say you wanted to add a constant to a register and store it in another? Well, at the moment we would only have 5-bits available to hold a number, that's not a lot. The max constant value would be $2^5 - 1 = 31$ (assuming unsigned number), that's not very big. Instead we have a different instruction format:

op	rs	rt	constant or address
6 bits	5 bits	5 bits	16 bits

Note that instructions with this format type are not just add immidiate type instructions, they are also used for load word instructions (among others). The following instruction

```

1 lw      $t0 , 1200($t1)    # $t0 <— A[300]
2 add    $t0 , $s2 , $t0      # $t0 <— h + A[300]
3 sw      $t0 , 1200($t1)    # A[300] <— $t0

```

is expressed as follows:

Op	rs	rt	rd	address/ shamt	funct
35	9	8		1200	
0	18	8	8	0	32
43	9	8		1200	

and in binary:

100011	01001	01000	0000	0100	1011	0000
000000	10010	01000	01000	00000	100000	
101011	01001	01000		0000	0100	1011 0000

From this, you can see that each register has an associated reference number:

NAME	NUMBER	USE	PRESERVED ACROSS A CALL?
\$zero	0	The Constant Value 0	N.A.
\$at	1	Assembler Temporary	No
\$v0-\$v1	2-3	Values for Function Results and Expression Evaluation	No
\$a0-\$a3	4-7	Arguments	No
\$t0-\$t7	8-15	Temporaries	No
\$s0-\$s7	16-23	Saved Temporaries	Yes
\$t8-\$t9	24-25	Temporaries	No
\$k0-\$k1	26-27	Reserved for OS Kernel	No
\$gp	28	Global Pointer	Yes
\$sp	29	Stack Pointer	Yes
\$fp	30	Frame Pointer	Yes
\$ra	31	Return Address	Yes

Here are the opcodes for the instructions we have seen so far:

instruction	format	op	rs	rt	rd	shamt	funct	address
add	R	0	reg	reg	reg	0	32 _{ten}	n.a.
sub (subtract)	R	0	reg	reg	reg	0	34 _{ten}	n.a.
add immediate	I	8 _{ten}	reg	reg	n.a.	n.a.	n.a.	constant
lw (load word)	I	35 _{ten}	reg	reg	n.a.	n.a.	n.a.	address
sw (store word)	I	43 _{ten}	reg	reg	n.a.	n.a.	n.a.	address

2.5.3 Check Yourself

What MIPS instruction does this represent?

op	rs	rt	rd	shamt	funct
0	8	9	10	0	34

From the opcode of 0 and function code of 34 we know this is a subtract function call.

We are dealing with registers 8, 9, 10 which are \$t0, \$t1, and \$t2 respectively. Because the MIPS register arguments are ordered rd, rs, rt we have get the MIPS code

```
1 sub      $t2 , $t0 , $t1
```

SUBSECTION 2.7

Instructions for Making Decisions

2.7.1 Branch Instructions

MIPS has two decision making instructions:

- **beq register1, register2, L1**
Go to the statement labeled L1 if the value in register1 equals the value in register2;
- **bne register1, register2, L1**
Go to statement labeled L1 if value in register1 does not equal value in register2.

For example, the C code

```
1 if ( i == j )
2     f = g + h;
3 else
4     f = g-h;
```

where f through j correspond to \$s0 through \$s4, corresponds to the MIPS code

```
1 bne    $s3 , $s4 , Else
2 add    $s0 , $s1 , $s2      # f = g + h
3 j      Exit                 # go to Exit
4
```

```

5 Else:
6     sub    $s0 , $s1 , $s2      # f = g-h
7
8 Exit:

```

Notice how we did things in reverse; instead of checking for equality (like in the C code), we checked for inequality.

2.7.2 Loops

How could we use our branch statements to compile the following while Loop:

```

1 while ( save[ i ] == k )
2     i += 1;

```

where `i` and `k` correspond to registers `$s3` and `$s5`. The address of `save[0]` is in `$s6`.

```

1 Loop:
2     sll    $t1 , $s3 , 2      # $t1 <-- i * 4
3     add    $t1 , $t1 , $s6      # $t1 <-- address of save[ i ]
4     lw     $t0 , 0($t1)      # $t0 <-- save[ i ]
5     bne    $t0 , $s5 , Exit   # if ( save[ i ] != k ) Exit
6     addi   $s3 , $s3 , 1      # i += 1
7     j      Loop              # go to Loop
8
9 Exit:

```

Because MIPS is byte addressed, to get the memory address of `save[i]` we need to multiply $i \cdot 4$ and add that to the base address. We then load that value to a temporary register, and test the condition.

2.7.3 Unsigned Comparison Tests

In addition to testing for equality, we might also want to test for less than, or greater than. To test if less than the instruction “set on less than”

```
1 sltu    $t0 , $s3 , $s4      # $t0 = 1 if $s3 < $s4; otherwise 0
```

This sets the destination register to 1 if less than, and 0 if not. An immediate version of this instruction also exists:

```
1 sltiu   $t0 , $s2 , 10      # $t0 = 1 if $s2 < 10, 0 otherwise
```

2.7.4 Signed Comparison Tests

We have to remember that with negative numbers, the most significant bit is turned on, and so using an unsigned comparison between -1 and 5 : $1111_{\text{TWO}} < 0101_{\text{TWO}}$ would return false. Thus, for signed comparisons we have `slt` and `slti`:

The assembler relieves the compiler and the assembly language programmer from the tedium of calculated addresses for branches.

This is an example of a **basic block**. This is a sequence without branches, except possibly at the end, and without branch targets or branch labels, except possibly at the beginning.

All relative conditions equal, not equal, less than, less than or equal, greater than, and greater than or equal, can all be created using `slt`, `slti`, `bne`, and `$zero` register.

```

1    slt      $t0 , $s3 , $s4      # $t0 = 1 if $s3 < $s4; otherwise 0
2    slti     $t0 , $s2 , 10      # $t0 = 1 if $s2 < 10, 0 otherwise

```

2.7.5 Bounds Checking Shortcut

Let's say we want to check if both

$$x < y \quad \text{and} \quad x \geq 0$$

where x is a and y are signed integers (and $y > 0$). Assume that $\$s0$ holds x and $\$s1$ holds y . Then we could do

```

1  # Check if x < y
2  slt      $t0 , $s0 , $s1
3  beq     $t0 , $zero , False
4
5  # Check if x >= 0
6  slt      $t0 , $s0 , $zero      # If x < 0
7  bne     $t0 , $zero , False
8
9  # Statement is true
10 j       True

```

We can actually check for both $x < y$ and $x \geq 0$ in one go by treating x and y as unsigned integers. The idea here is that since a negative number will have its most significant bit turned on, it looks like a very large unsigned number. Therefore no matter how large y is, if x is a negative number it will appear as a larger unsigned number. Thus,

$$0 \leq x \leq y$$

will only be true if x is nonnegative and less than y , and when treated as unsigned numbers, testing if $x < y$ will also check if $x \geq 0$:

```

1  sltu $t0 , $s0 , $s1
2  beq   $t0 , $zero , False
3  j     True

```

This shortcut is really helpful if we want to check if some index is out of bounds for a given array. For example, let's say we have an array `example` whose upper bounds y . Then we could use the code above to check if some index i is in bounds ($0 \leq i < y$).

SUBSECTION 2.8

Supporting Procedures in Computer Hardware

Thus far, we have only really learned how to perform a sequence of instructions, but what if we want to perform that same sequence of instructions a lot of times, and maybe we want to act upon some variables. In C we have functions, and in assembly, we have **procedures**. To perform a procedure in MIPS, we follow the following steps:

1. Put parameters in a place where the procedure can access them.
2. Transfer control to the procedure.

A **procedure** is a stored subroutine that performs a specific task based on the parameters with which it is provided.

3. Acquire the storage resources needed for the procedure.
4. Perform the desired task.
5. Put the result value in a place where the calling program can access it.
6. Return control to the point of origin, since a procedure can be called from several points in a program.

The following registers are used for procedure calling

- \$a0-\$a3: four argument registers in which to pass parameters
- \$v0-\$v1: two value registers in which to return values
- \$ra: one return address register to return to the point of origin

The return address register is populated by the **jump-and-link-instruction** (jal):

jal ProcedureAddress. This instruction simultaneously jumps to ProcedureAddress, and populates the return address register (\$ra) with the address of the following instruction. This way the computer knows what to perform once the procedure is finished.

The **caller** is the one that puts the parameter values in \$a0-\$a3 and uses jal X to jump to procedure X (**callee**). Once the callee has finished, it places the results in \$v0 and \$v1, and then returns control to the caller using jr \$ra. This is an unconditional jump to the address specified in a register.

To make all of this work, we need to store the address of the current instruction being executed. We do this in a register called the **program counter** (PC). The jal instruction saves PC + 4 in register \$ra.

The **return address** is a link to the calling site that allows a procedure to return to the proper address; in MIPS it is stored in register \$ra.

2.8.1 Nested Procedures

Let's say you have the following program:

```
1 #include <stdio.h>
2
3 int fact(int n)
4 {
5     if (n < 1)
6         return 1;
7     else
8         return (n * fact(n-1));
9 }
10
11 int main() {
12     int result = fact(3);
13     printf("%d\n", result);
14     return 0;
15 }
```

Because **fact(n)** can call itself, you have to be careful about saving the return address register (\$ra), as well as any other registers that may be used (more info below). We can write the above program in MIPS as follows:

```

1 main:
2 addi    $a0, $zero, 3
3 jal     fact
4 add    $a0, $v0, $zero # store result for printing
5 addi    $v0, $zero, 1  # print int syscall
6 syscall
7 addi    $v0, $zero, $10 # exit syscall
8 syscall
9
10 fact:
11 addi   $sp, $sp, -8   # make space for two items on stack
12 sw     $ra, 4($sp)   # save return address
13 sw     $a0, 0($sp)   # save arg n
14 slti   $t0, $a0, 1   # test if n < 1
15 beq    $t0, $zero, L1 # if n >= 1, goto L1
16 lw     $ra, 4($sp)   # restore $ra
17 lw     $a0, 0($sp)   # restore arg n
18 addi   $v0, $zero, 1  # return 1
19 addi   $sp, $sp, 8   # pop two items from stack
20 jr     $ra             # return to caller
21
22 L1:
23 addi   $a0, $a0, -1   # n += -1
24 jal    fact           # calls fact and $ra <— PC+4
25 lw     $a0, 0($sp)   # restore arg n
26 lw     $ra, 4($sp)   # restore $ra
27 addi   $sp, $sp, 8   # Pops two items from stack
28 mul    $v0, $a0, $v0   # return n*fact(n-1)
29 jr     $ra             # return to caller

```

Let's look at how the register values, and stack changes as we go through the program:

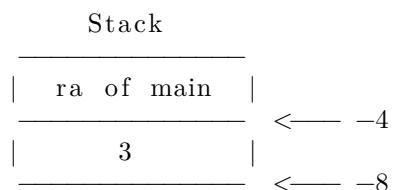
1. In **main** (before first call of **fact**):

```

$a0 <— 3
$ra <— return to main

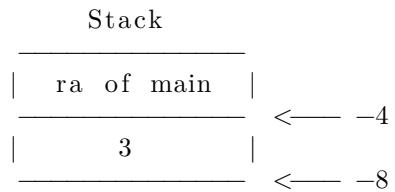
```

2. During first call of **fact(3)**



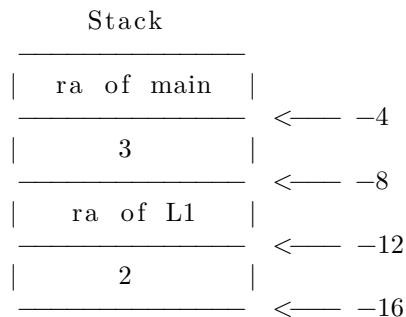
3. First break to L1:

$\$a0 \leftarrow 3-1 = 2$
 $\$ra \leftarrow \text{return to L1}$



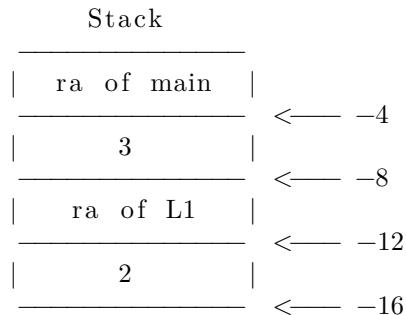
4. During second call to **fact(2)**

$\$a0 = 3-1 = 2$
 $\$ra = \text{return to L1}$



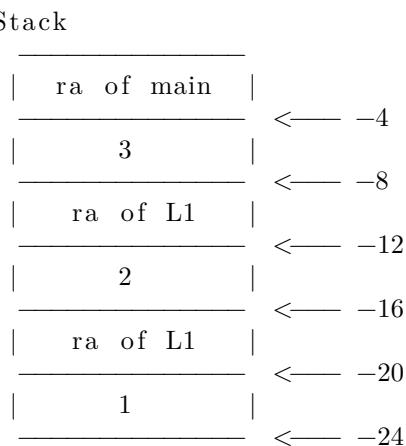
5. During second break to L1

$\$a0 \leftarrow 2-1 = 1$
 $\$ra \leftarrow \text{return to L1}$



6. During third call to **fact(1)**

$\$a0 = 1$
 $\$ra = \text{return to L1}$



7. During third break to L1

$\$a0 <-- 1 - 1 = 0$ $\$ra <-- L1$	Stack <hr/> <table border="0" style="width: 100%; border-collapse: collapse;"> <tr> <td style="padding: 0 10px;"> ra of main </td><td style="padding: 0 10px; text-align: right;">← -4</td></tr> <tr> <td style="padding: 0 10px;"> 3 </td><td style="padding: 0 10px; text-align: right;">← -8</td></tr> <tr> <td style="padding: 0 10px;"> ra of L1 </td><td style="padding: 0 10px; text-align: right;">← -12</td></tr> <tr> <td style="padding: 0 10px;"> 2 </td><td style="padding: 0 10px; text-align: right;">← -16</td></tr> <tr> <td style="padding: 0 10px;"> ra of L1 </td><td style="padding: 0 10px; text-align: right;">← -20</td></tr> <tr> <td style="padding: 0 10px;"> 1 </td><td style="padding: 0 10px; text-align: right;">← -24</td></tr> </table> <hr/>	ra of main	← -4	3	← -8	ra of L1	← -12	2	← -16	ra of L1	← -20	1	← -24
ra of main	← -4												
3	← -8												
ra of L1	← -12												
2	← -16												
ra of L1	← -20												
1	← -24												

8. During fourth call to **fact(0)**

$\$a0 = 0$ $\$ra = L1$ $\$ra <-- \text{stack}[-28] = L1$ $\$a0 <-- \text{stack}[-32] = 0$ $\$v0 <-- 1 * 1 = 1$	Stack <hr/> <table border="0" style="width: 100%; border-collapse: collapse;"> <tr> <td style="padding: 0 10px;"> ra of main </td><td style="padding: 0 10px; text-align: right;">← -4</td></tr> <tr> <td style="padding: 0 10px;"> 3 </td><td style="padding: 0 10px; text-align: right;">← -8</td></tr> <tr> <td style="padding: 0 10px;"> ra of L1 </td><td style="padding: 0 10px; text-align: right;">← -12</td></tr> <tr> <td style="padding: 0 10px;"> 2 </td><td style="padding: 0 10px; text-align: right;">← -16</td></tr> <tr> <td style="padding: 0 10px;"> ra of L1 </td><td style="padding: 0 10px; text-align: right;">← -20</td></tr> <tr> <td style="padding: 0 10px;"> 1 </td><td style="padding: 0 10px; text-align: right;">← -24</td></tr> <tr> <td style="padding: 0 10px;"> ra of L1 </td><td style="padding: 0 10px; text-align: right;">← -28</td></tr> <tr> <td style="padding: 0 10px;"> 0 </td><td style="padding: 0 10px; text-align: right;">← -32\$</td></tr> </table> <hr/>	ra of main	← -4	3	← -8	ra of L1	← -12	2	← -16	ra of L1	← -20	1	← -24	ra of L1	← -28	0	← -32\$
ra of main	← -4																
3	← -8																
ra of L1	← -12																
2	← -16																
ra of L1	← -20																
1	← -24																
ra of L1	← -28																
0	← -32\$																

9. Backtracking to 3rd L1:

$\$ra \leftarrow \text{stack}[-20] = \text{L1}$ $\$a0 \leftarrow \text{stack}[-24] = 1$ $\$v0 \leftarrow 1 * 1 = 1$	**Stack** ---							--	--	--	--	--																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																	

- The first instruction in the function manipulates `$sp` to allocate F words of stack:

```
addi $sp, $sp, -F*4
```

where $F = A + L + P + (1 \text{ for } \$ra \text{ if needed}) + (1 \text{ for } \$fp \text{ if in use}) + (2 \text{ for } \$gp) + \text{padding}$. Also, F is the frame size in words, and should be even.

Corollary 1: `$sp` is not changed at any other point than the first and last instructions in the function.

- If you want to use any preserved register in any place during the function, you must store them into the stack when you **enter** the function (not when you use it). Let's assume we are saving `$s0`, `$s1`, and `$s2` registers, as well as `$ra` and `$sp`. Then we would write:

```
sw $ra, (F-1)*4($sp)    # Not needed for leaf functions
sw $fp, (F-2)*4($sp)
sw $s2, (F-3)*4($sp)
sw $s1, (F-4)*4($sp)
sw $s0, (F-5)*4($sp)
```

- After saving preserved registers, copy `$sp` into `$fp`, so that from now until you start returning, you can use offsets from `$fp` to access the locations. Save and restore local variables as needed after this point, using `$fp` and the spots already allocated for them on the stack.
- You may pass information **into** a function only through `$a0-$a3` and extra stack arguments. As the called function, you may not use the value held in any other register except `$sp`, `$fp`, `$gp` and `$ra` (i.e. don't try and use `$v0` or `$s0`). The 4 (fifth) argument to the current function is available at $(F + 4) * 4(\$fp)$. You may store argument 0 from `$a0` into its reserved spot at $F * 4(\$fp)$, arg1 into $(F + 1) * 4(\$fp)$, etc., if you want.
- You may pass information **out of** a function through only `$v0-$v1`.
- You may only enter a function at its beginning. You may not **jal** into the middle of a function.
- You may have only one `jr $ra` in each function. It is the last instruction in the function, following the instructions that restore the registers saved.

Extra rules for non-leaf functions:

- Look at the declarations for all procedures this function *might* call. Find the one with the largest number of arguments, A , and then make A at least 4. You must allocate $A \cdot 4$ bytes on the stack upon entry into your function.
- If passing more than 4 arguments, then place 4 arguments into `$a0-$a3`. Every argument past the 4th argument is placed on the stack. For example suppose 6 arguments are in `$s0-$s5`:

```
move $a0, $s0    # 0($fp) reserved on the stack
move $a1, $s1    # 4($fp) reserved on the stack
move $a2, $s2    # 8($fp) reserved on the stack
move $a3, $s3    # 12($fp) reserved on the stack
sw $s4, 16($fp) # 5th arg in the stack. Always here!
sw $s5, 20($fp) # 5th arg in the stack. Always here!
```

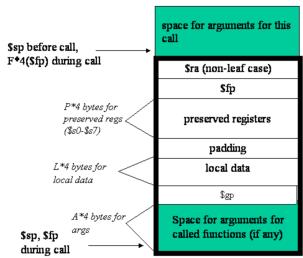


Figure 9. Stack organization during call

- You must preserve register \$ra because a called-function jal destroys the value, and it is a preserved register.

SUBSECTION 2.10

MIPS Addressing for 32-bit Immediates and Addresses

2.10.1 32-bit Constant into Register

If you have been following closely, you may have noticed that when adding a constant to a register there are only 16 bits available for the constant (as it is an I-type instruction):

op	rt	rs	constant
6	5	5	16

The question then remains, what do you do if you want to place a constant larger than 16 bits? You achieve this with the following two instructions:

1. **lui rt, immediate**
which loads immediate value into the upper 16 bits of **rt** and zeros lower 16 bits.
2. **ori rt, rs, immediate**
Bitwise or of **rs** and **immediate** placed into **rt** ($rt \leftarrow rs \mid imm$).

A word of caution. Using the instruction **addi** sign extends the msb of immediate.

2.10.2 Branch Addressing

Branch instructions take are I-type instructions:

op	rs	rt	encoded address
6	5	5	16

The target of the encoded address is

$$\text{Target} = 4(\text{encoded address}) + \text{PC} + 4.$$

That is, branch instructions are word addressed ($4(\text{encoded address})$), and relative to the next instruction ($\text{PC} + 4$).

2.10.3 Jump Addressing

Jump instructions are J-type instructions:

op	encoded address
6	26

In this case the lower 28 bits of the address to be jumped to are

$$4 \cdot \text{encoded address}$$

and the upper four are that of the PC. If you need to jump to a really far away address (more than 64 million instructions away), then you use jump register instruction (**jr**).

SUBSECTION 2.12

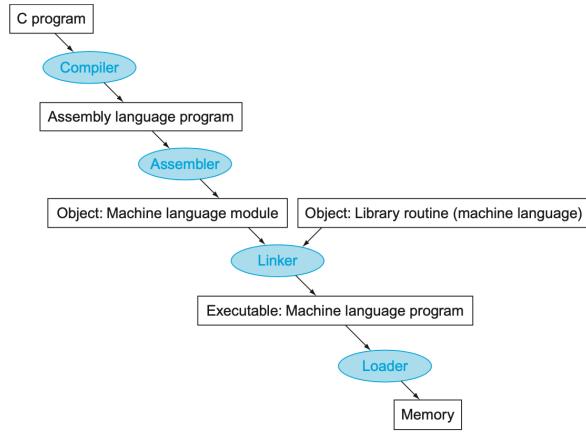
Translating and Starting a Program

Figure 10. The four steps that go into loading your C program into memory.

2.12.1 Compiler

The compiler transforms the C program into an assembly language program.

2.12.2 Assembler

The main job of the assembler is to convert assembly language into binary machine code (object file). The assembler must determine the addresses associated with labels, which it keeps track of in a symbol table which contains pairs of symbols and addresses. The object for UNIX systems typically contains six distinct pieces:

- The *object file header* describes the size and position of the other pieces of the object file.
- The *text segment* contains the machine language code.
- The *static data segment* contains data allocated for the life of the program. (UNIX allows programs to use both static data, which is allocated throughout the program, and dynamic data, which can grow or shrink as needed by the program.)
- The *relocation information* identifies instructions and data words that depend on absolute addresses when the program is loaded into memory.
- The *symbol table* contains the remaining labels that are not defined, such as external references.
- The *debugging information* contains concise descriptions of how the modules were compiled so that a debugger can associate machine instructions with C source files and make data structures readable.

2.12.3 Linker

The linker or link editor is a systems program that combines independently assembled machine language programs and resolves all undefined labels into an executable file. In this way, you can combine your program with other programs (such as header files or libraries) without having to recompile everything together.

2.12.4 Loader

The loader places an object program in main memory so that it is ready to execute. The loader follows the following steps in UNIX systems:

1. Reads the executable file header to determine size of the text and data segments.
2. Creates an address space large enough for the text and data.
3. Copies the instructions and data from the executable file into memory.
4. Copies the parameters (if any) to the main program onto the stack.
5. Initializes the machine registers and sets the stack pointer to the first free location.
6. Jumps to a start-up routine that copies the parameters into the argument registers and calls the main routine of the program. When the main routine returns, the start-up routine terminates the program with an `exit` system call.

SUBSECTION 2.14

Arrays versus Pointers

Let's look at two ways we can set an array to all zeros: one using pointers, and one using arrays, so that we get an idea for how C pointers work.

2.14.1 Array

Let's say that we have the following C code:

```

1 clear1( int array [] , int size )
2 {
3     int i ;
4     for ( i = 0; i < size ; i +=1)
5         array [ i ] = 0;
6 }
```

Then we would have the following MIPS code:

```

1 addiu $sp, $sp, $-16      # Space for four arguments
2 addi $sp, $sp, $-32      # Space for array of 8 words
3 addi $sp, $sp, $-4       # Space for $gp
4 addi $sp $sp, $-4        # space for $fp
5 addi $sp $sp $-4        # Space for $ra
6
7 F = 4(arguments) + 8(local data) + 1(ra) + 1(fp) + 2(gp)
8 = 16
9 Stack
10 _____
11 |    ra    |
12 |_____| <-- +60
13 |    fp    |
14 |_____| <-- +56
15 |    array [7] |
16 |_____| <-- +52
17 |    array [6] |
```

```

18      _____ <--- +48
19 |   array [5] |
20 | _____ <--- +44
21 |   array [4] |
22 | _____ <--- +40
23 |   array [3] |
24 | _____ <--- +36
25 |   array [2] |
26 | _____ <--- +32
27 |   array [1] |
28 | _____ <--- +28
29 |   array [0] |
30 | _____ <--- +24
31 |
32 | _____ <--- +20
33 |   gp   |
34 | _____ <--- +16
35 |   arg 3 spce |
36 | _____ <--- +12
37 |   arg 2 spce |
38 | _____ <--- +8
39 |   arg 1 spce |
40 | _____ <--- +4
41 |   arg 0 spce |
42 | _____ <--- +0
43
44 main:
45 addiu $sp, $sp, -64          # Allocate stack space
46 sw    $ra, 60($sp)           # Save return address
47 sw    $fp, 56($sp)           # Save frame pointer
48 sw    $gp, 16($sp)           # Save global pointer
49 add   $fp, $sp, $zero        # Establish frame pointer
50 addiu $a0, $fp, 24           # arg0 <--- &array[0]
51 addiu $a1, $zero, 8          # arg1 <--- size = 8
52 jal   clear1                # clear1(&array, 8)
53 j exitclear1: add t0,zero, zeroi <-- 0slt t3, a1, 1t3 <- 1 if size < 1 bne t3, zero,
exit  if (size < 1) exitloop: sll t1,t0, 2 t1 <-- i * 4addt2, a0,t1 t2 <-- memaddressofarray[i] swzero, 0(t2)array[i] <-- 0addit0, t0, 1i+ = 1slt t3, t0,a1
t3 <-- i, if i < size bnet3, zero, Loopgotoloopifiexit :

```

SECTION 3

Arithmetic for Computers

SUBSECTION 3.2

Addition and Subtraction

Addition in binary works just like addition in decimal, going bit by bit, and carrying over:

$$\begin{array}{r}
10_{\{10\}} = 01010_2 \\
+ 6_{\{10\}} = 00110_2
\end{array}$$

Note that this code only works if `size > 0`; ANSI C requires a test of size before the loop, but we'll skip that legality here.

$$16_{\{10\}} = 10000_{\{2\}}$$

Subtraction can either be done via the subtract operation, or by adding a negative number (in two's complement):

$$\begin{array}{r} 0000\ 0000\ 0000\ 0000\ 0000\ 0000\ 0000\ 0111_{\text{two}} = 7_{\text{ten}} \\ - 0000\ 0000\ 0000\ 0000\ 0000\ 0000\ 0000\ 0110_{\text{two}} = 6_{\text{ten}} \\ \hline = 0000\ 0000\ 0000\ 0000\ 0000\ 0000\ 0000\ 0001_{\text{two}} = 1_{\text{ten}} \end{array}$$

or via addition using the two's complement representation of -6 :

$$\begin{array}{r} 0000\ 0000\ 0000\ 0000\ 0000\ 0000\ 0000\ 0111_{\text{two}} = 7_{\text{ten}} \\ + 1111\ 1111\ 1111\ 1111\ 1111\ 1111\ 1111\ 1010_{\text{two}} = -6_{\text{ten}} \\ \hline = 0000\ 0000\ 0000\ 0000\ 0000\ 0000\ 0000\ 0001_{\text{two}} = 1_{\text{ten}} \end{array}$$

3.2.1 Overflow Twos Complement

Addition

- Cannot occur when adding operands with different signs. This is because the sum can be no larger than one of the numbers, and both numbers fit.
- Overflow can be detected when adding numbers of the same sign but the sign bit changes.

Subtraction

- Cannot occur when signs are the same ($a - b = a + (-b)$).
- Overflow occurs when we subtract a negative number from a positive number but get a negative number ($a - (-b) = -x$)

3.2.2 Unsigned Overflow

Depending upon the situation, one may or may not want an exception to be thrown when overflow occurs. For this reason there are two types of integer operations.

- **add**, **addi**, and **sub** cause exceptions on overflow.
- **addu**, **addiu**, and **subu** do not cause exceptions on overflow.

Note that C ignores overflows, and so MIPS compilers will always generate the unsigned versions.

```
addu $t0, $t1, $t2 # $t0 = sum, but don't trap
xor $t3, $t1, $t2 # Check if signs differ
slt $t3, $t3, $zero # $t3 = 1 if signs differ
bne $t3, $zero, No_overflow # $t1, $t2 signs *,
                           # so no overflow
xor $t3, $t0, $t1 # signs =; sign of sum match too?
                   # $t3 negative if sum sign different
slt $t3, $t3, $zero # $t3 = 1 if sum sign different
bne $t3, $zero, Overflow # All 3 signs !=; goto overflow
```

For unsigned addition ($\$t0 = \$t1 + \$t2$), the test is

```
addu $t0, $t1, $t2      # $t0 = sum
nor $t3, $t1, $zero     # $t3 = NOT $t1
                       # (2's comp - 1:  $2^{32} - \$t1 - 1$ )
slt $t3, $t3, $t2       # ( $2^{32} - \$t1 - 1$ ) < $t2
                       #  $\Rightarrow 2^{32} - 1 < \$t1 + \$t2$ 
bne $t3,$zero,Overflow # if( $2^{32}-1 < \$t1+\$t2$ ) goto overflow
```

Note that although **addiu** is unsigned addition, the 16-bit immediate field is sign extended to 32-bits and so the immediate field is signed, even if the operation is “unsigned”.

SUBSECTION 3.3

Multiplication

Let's say we want to multiply 1011×0101 , then we would perform the following steps:

$$\begin{array}{r}
 1011 \quad \textit{Multiplicand} \\
 \times 0101 \quad \textit{Multiplier} \\
 \hline
 1011 \\
 0000 \\
 1011 \\
 + 0000 \\
 \hline
 00110111 \quad \textit{Product}
 \end{array}$$

Notice that we just add copies of the multiplicand (shifted over by 1 each time), or zero (depending upon bit of multiplier). Thus, to perform multiplication, we will need the following pieces of hardware:

- A multiplicand register. Because we are shifting over the multiplicand each time, it needs to be twice as big (so 64-bits).
- A multiplier register. Size of 32-bits.
- A product register. Needs to be 64 bits wide.
- An Arithmetic Logic Unit (ALU) to perform the addition of product and multiplicand registers.
- A control test that in each cycle tells:
 - the ALU to sum the multiplicand and product (if control says to)
 - the multiplicand to shift left,
 - the multiplier to shift right,
 - the product register to change or not depending upon the lsb if the multiplier register.

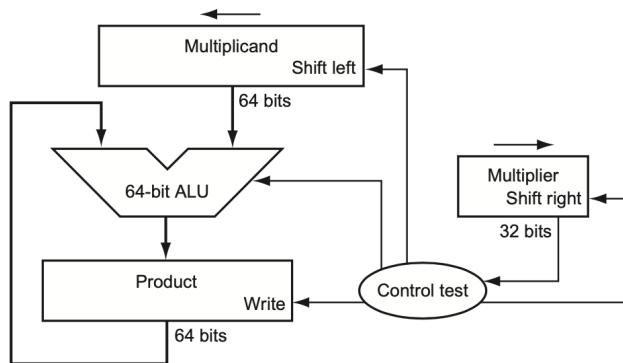


Figure 11. A simple version of multiplication hardware.

So what will the first cycle of hardware look like?

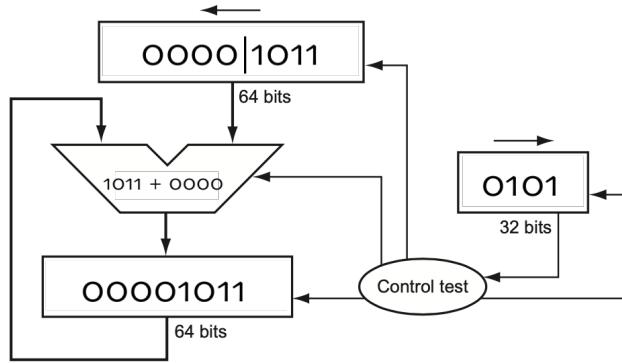


Figure 12. So at first, the product register is set to 0000, the multiplicand register is set to 1011, and the multiplier register is set to 0101. The control test tests the first bit of the multiplicand (1), and so tells the ALU to sum the product and multiplicand registers, and store the value in the product register. The control then shifts the multiplicand 1 bit to the left, and the multiplier 1 bit to the right.

3.3.1 Faster Multiplication

Parallel Processing Speedup

- The multiplier and multiplicand are shifted while the multiplicand is added to the product if the multiplier bit is a 1.
- The hardware just has to ensure that it tests the right bit of the multiplier and gets the preshifted version of the multiplicand.

Resource Optimization Speedup If you look back at doing the multiplication by hand, you may notice that at each step you are really only adding four bits at once. What we will do now, is have a 32-bit multiplicand register, a 32-bit adder, and a 64-bit product, and shift the product register instead of the multiplicand register. Let's see how this works in practice.

- Step 1**
Multiplicand: 1011
Multiplier: 0101
Product: 1011 0000
- Step 2**
Multiplicand: 1011
Multiplier: 0101
Product: 0101 1000
- Step 3**
Multiplicand: 1011
Multiplier: 0101
Product: 1101 1100
- Step 4**
Multiplicand: 1011
Multiplier: 0101
Product: 0110 1110

The last shift will then set the product as 0011 0111.

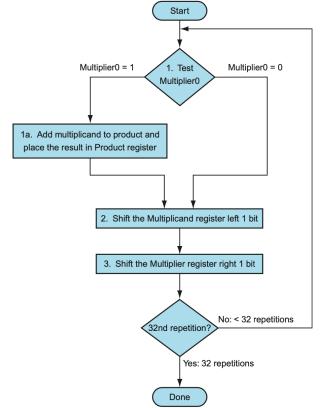


Figure 13. The multiplication algorithm.

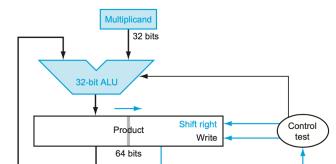


Figure 14. The faster multiplication hardware, taking advantage of parallel processing as well as reduced register size

3.3.2 Even Faster Multiplication

SUBSECTION 3.4

Division

Let's first examine a little bit more closely the long division from elementary school:

$$\begin{array}{r} 1001 \\ 1000 \overline{)1001010} . \\ \underline{1000} \\ 1010 \\ \underline{1000} \\ 10 \end{array}$$

Notice how as we are testing to see the divisor is less than the n MSBs of the dividend where n increases by 1 each cycle. If the divisor is less than, we subtract the divisor from the dividend and put 1 in the quotient.

So how does all of this transfer over to hardware?

- Well, so that we are comparing divisor to the MSBs of the dividend we will store the 32-bit divisor in a 64-bit integer, but start with it in the 32 MSBs. This way, after each cycle we just shift the divisor to the right one.
- Because we keep subtracting from the dividend, we just make our life simple and start the dividend as the remainder.
- Instead of placing the 1 or 0 of the quotient starting from the LSB, we place start with the quotient as zero and shift the quotient left 1 and set LSB to 1 if applicable. This way at the end this will get shifted over to become the MSB.

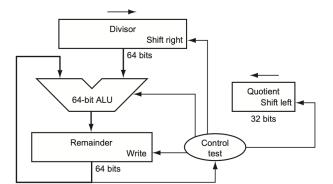


Figure 15. Hardware implementation of division

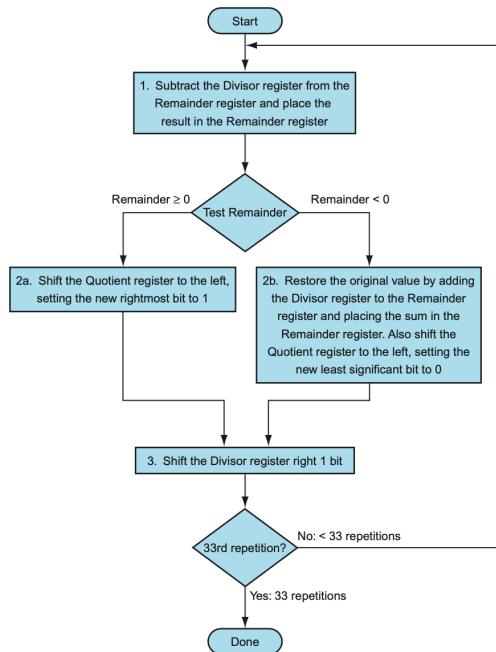


Figure 16. The division algorithm that we will implement in hardware.

3.4.1 Faster Division

We can make this a little bit more efficient. Instead of shifting the divisor right, we can shift the remainder/quotient left. So we start with the quotient in the 32 LSBs, shift the remainder left one cycle and compare the 32 MSBs of the remainder with the divisor (we want the divisor to be less than the upper 32 MSBs of the quotient/remainder). In this way we now only need a 32 bit ALU.

At the end of this you may notice that the 32 LSBs of the quotient/remainder are all zeros. To utilize this space we will use this as the spot for the remainder to go, saving us an extra 32 bits.

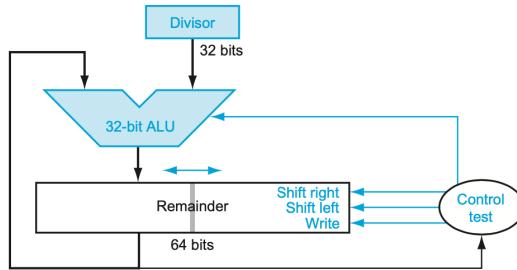


Figure 17. The faster division hardware

SECTION 4

Basics of Logic Design

For the most part we will be focusing on *combinational logic*. This is logic that has no memory (state) and therefore

$$\text{Same Input} \Leftrightarrow \text{Same Output} .$$

SUBSECTION 4.1

Truth Tables

Because of the above fact, a combinational block can be completely described by a truth table. In a logic block an input is either true or false. Therefore, a logic block with n inputs has 2^n possible different possibilities.

Let's consider the logic function with inputs (A, B) and outputs (C, D, E) where:

- C is true if no inputs are true,
- D is true if at least one input is true, and
- E is true if all inputs are true.

The truth table would be defined as follows

Inputs		Outputs		
A	B	C	D	E
0	0	1	0	0
0	1	0	1	1
1	0	0	1	0
1	1	0	1	1

Sometimes truth tables are hard to read, so we can express each output of the truth table as a boolean expression:

$$\begin{aligned} C &= \bar{A} \cdot \bar{B} \\ D &= A + B \\ E &= A \cdot B. \end{aligned}$$

SUBSECTION 4.2

Boolean Algebra

There are three elementary operators:

- logical OR: $A + B$,
- logical AND: $A \cdot B$, and
- logical NOT: \bar{A} .

4.2.1 Laws

Identity Laws: $A + 0 = A$

$$A \cdot 1 = A$$

Zero and One Laws: $A + 1 = 1$

$$A \cdot 0 = 0$$

Inverse Laws: $A + \bar{A} = A$

$$A \cdot \bar{A} = 0$$

Commutative Laws: $A + B = B + A$

$$A \cdot B = B \cdot A$$

Associative Laws: $A + (B + C) = (A + B) + C$

$$A \cdot (B \cdot C) = (A \cdot B) \cdot C$$

Distributive Laws: $x \cdot (y + z) = (x \cdot y) + (x \cdot z)$

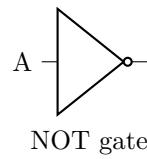
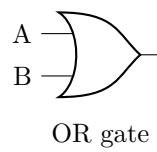
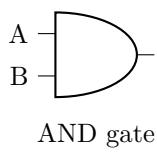
DeMorgan's Laws: $\overline{(A + B)} \Leftrightarrow \bar{A} \cdot \bar{B}$

$$\overline{(A \cdot B)} \Leftrightarrow \bar{A} + \bar{B}.$$

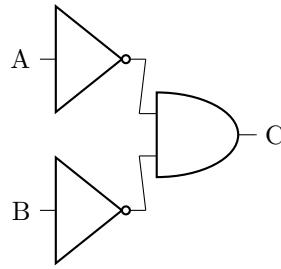
SUBSECTION 4.3

Logic Blocks

We can implement truth tables and boolean expressions with the use of three types of logic gates:



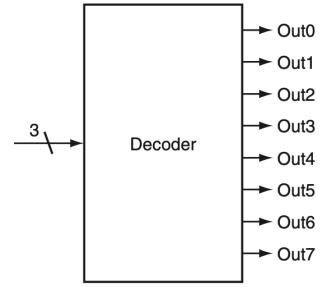
For example, we can build function $C = \bar{A} \cdot \bar{B}$ as follows:



4.3.1 Decoder

A decoder is a logic block that has n -bit input and 2^n outputs, where only one output is asserted for each input combination. For example:

Inputs		Outputs			
A	B	C	D	E	F
0	0	0	0	0	1
0	1	0	0	1	0
1	0	0	1	0	0
1	1	1	0	0	0



4.3.2 Multiplexer

A multiplexer selects one of the given inputs based upon the select signal given. A multiplexer has three parts:

1. A decoder that generates n signals (from $\log_2 n$ inputs), each indicating a different input value
2. An array of n AND gates, each combining one of the inputs with a signal from the decoder
3. A single large OR gate that incorporates the outputs of the AND gates.

Figure 18. A 3 bit decoder

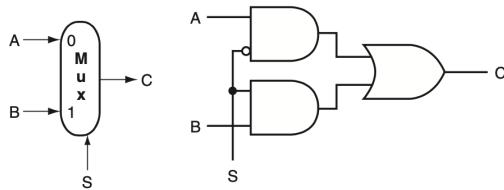


Figure 19. We can see either A or B is selected based upon if the value of S is 0 or 1.

The truth table for the above MUX will be

index	S	C
0	0	A
1	1	B

and in boolean algebra:

$$C = (\bar{S} \cdot A) + (S \cdot B).$$

4.3.3 Two-Level Logic

Any logic function can be written where every **input** is either true or a complemented variable and there are only two levels of gates:

- an AND level and
- an OR level.

There may be a possible inversion of the final output. This is a lot to unpack, so let's look at an example.

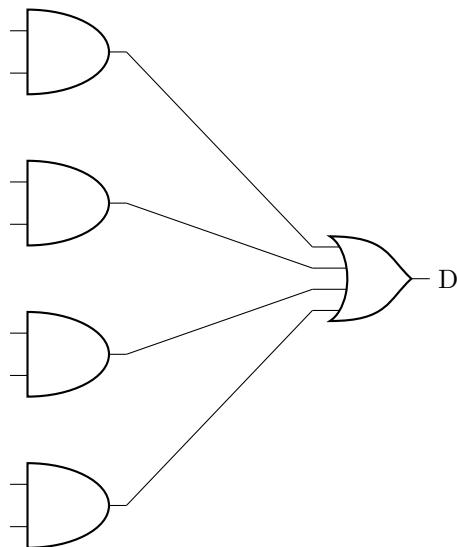
Let's say we are given the following truth table for D :

Inputs			Outputs
A	B	C	D
0	0	0	0
0	0	1	1
0	1	0	1
0	1	1	0
1	0	0	1
1	0	1	0
1	1	0	0
1	1	1	1

Notice that there are four input combinations for which D is true:

- $\bar{A} \cdot \bar{B} \cdot C$
- $\bar{A} \cdot B \cdot \bar{C}$
- $A \cdot \bar{B} \cdot \bar{C}$
- $A \cdot B \cdot C$

We call each of these terms a **product term**. For example the function D can be implemented in hardware using two-level logic as follows:



A **product term** or a **minterm** is a set of logical inputs joined by conjunction (AND operations); the product terms for the first logic stage of the PLA.

4.3.4 Programmable Logic Array

When dealing with a set of logic functions (multiple inputs and multiple outputs) there will be an

- AND gate for each unique set of inputs for which there is at least one corresponding true output
- OR gate for each output function.

This way of structured-logic implementation is called a programmable logic array. Let's say we have the following logic table:

Inputs			Outputs		
A	B	C	D	E	F
0	0	0	0	0	0
0	0	1	1	0	0
0	1	0	1	0	0
0	1	1	1	1	0
1	0	0	1	0	0
1	0	1	1	1	0
1	1	0	1	1	0
1	1	1	1	0	1

Then the associated PLA that implements the logic function is:

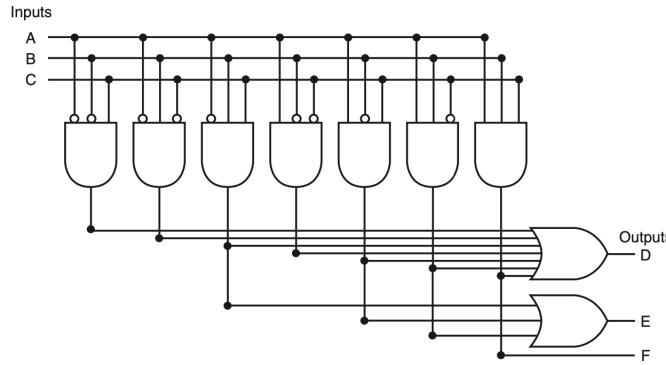


Figure 20. Notice that only 7 AND gates are required because no function is true for $A = B = C = 0$, and so we can ignore that. Also note that although there are three functions, there are only two OR gates because only combination of inputs turns F on.

4.3.5 Don't Cares

Don't cares, as the name suggests, are values of an input, or values of an output for which we don't care what the value is; the value could be either 0 or 1 and the effect of the logic block would be unchanged.

- Input Don't Cares: Arise when an output depends on only some of the inputs
- Output Don't Cares: Arise when we don't care about the value of an output because another output is true.

Don't cares are shown as X's on a truth table. They are important because they make it easier to implement the optimization of a logic function. For example, consider the following logic functions:

- D is true if A or C is true

- E is true if A or B is true
- F is true if exactly one of the inputs is true. But we don't care about F if both D and E are true

The truth table is as follows:

Inputs			Outputs		
A	B	C	D	E	F
0	0	0	0	0	0
0	0	1	1	0	1
0	1	0	0	1	1
0	1	1	1	1	0
1	0	0	1	1	1
1	0	1	1	1	0
1	1	0	1	1	0
1	1	1	1	1	0

Inputs			Outputs		
A	B	C	D	E	F
0	0	0	0	0	0
0	0	1	1	0	1
0	1	0	0	1	1
X	1	1	1	1	X
1	X	X	1	1	X

Figure 21. Notice that with the use of don't cares, we can have the truth table be smaller because of repeated rows.

4.3.6 Arrays of Logic Units

Up until now we have been performing combinational operations on just single bits, but we often want to do the same operations on whole words (32-bits). For example you may want to see if the value in one register is the same as the value as the other. Or you may want to have a MUX that selects between two words, and not just two bits:

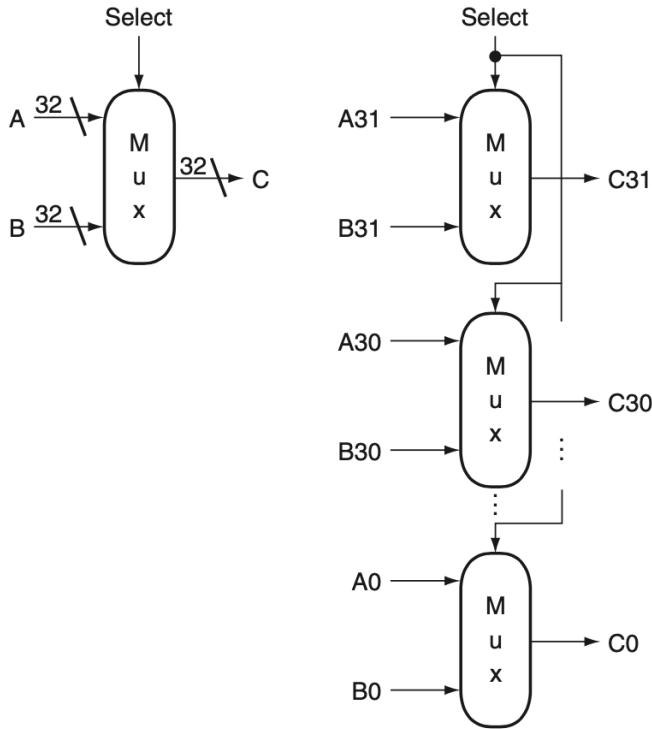


Figure 22. The 32-bit MUX is just 32 1-bit MUX's connected together.

SECTION 5

Constructing a Basic Arithmetic Logic Unit

An ALU is the part of the processor that is responsible for performing both logical and arithmetic operations such as addition, subtraction, AND, and OR.

SUBSECTION 5.1

1-Bit ALU

Let's go about implementing each of the different operations an ALU can perform.

5.1.1 AND / OR Logical Operation

To implement this you just have a MUX that selects between the AND, and OR operations:

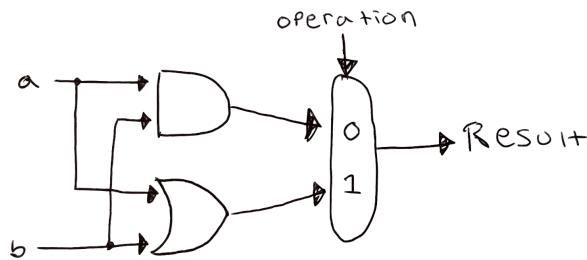


Figure 23. The multiplexor selects between and or or depending upon the opcode received.

5.1.2 Addition

To implement addition, with each bit in addition to looking at the bits adding together we also want to see if there was carryover from the the prior (less significant) bit. Thus, to deploy our operation in hardware we will need three **inputs**:

- a
- b
- carry in

and two **outputs**:

- computed sum
- carry out.

How is the carry out computed? Well, there is carry out if 2 or 3 of the inputs are true. Thus we have

$$\begin{aligned} \text{Carry Out} &= (a \cdot b) + (a \cdot \text{carry in}) + (b \cdot \text{carry in}) + (a \cdot b \cdot \text{carry in}) \\ &= (a \cdot b) + (a \cdot \text{carry in}) + (b \cdot \text{carry in}). \end{aligned}$$

Implementing this using logic gates we have:

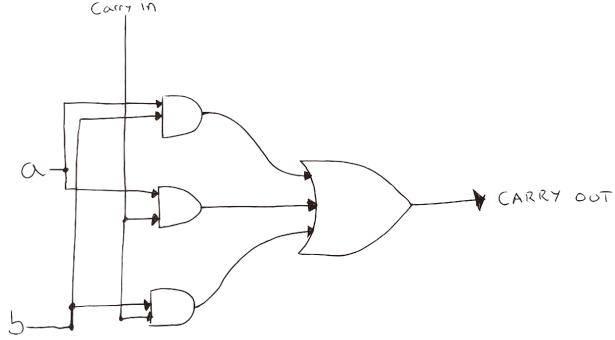


Figure 24. For cary out to be true we need to have t least two of the inputs to be true.

To calculate the **sum** using logic we see that sum will be true if an odd number of inputs are true (1 or 3). Therefore:

$$\text{sum} = (a \cdot \bar{b} \cdot \overline{\text{carry in}}) + (\bar{a} \cdot b \cdot \overline{\text{carry in}}) + (\bar{a} \cdot \bar{b} \cdot \text{carry in}) + (a \cdot b \cdot \text{carry in}).$$

Implementing this using logic gates we have:

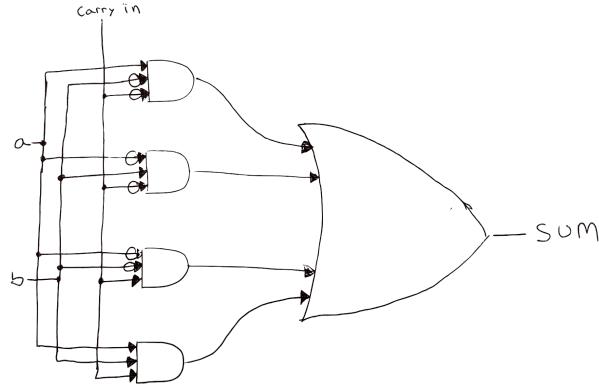


Figure 25. Sum is true if there are an odd number of true inputs.

5.1.3 Completing the 1-Bit ALU

To complete the 1-bit ALU we just connect the AND/OR/ADD operations to a mux that selects the desired operation.

SUBSECTION 5.2

32-Bit ALU

The 32-bit ALU is created by connecting 32 1-bit ALU together. Although we have our 1-bit ALU, there are some additional features that we want to make sure we have.

5.2.1 Subtraction

To compute subtraction we will use the addition we have already implemented and add a negative number. We can do this because $-b = \bar{b} + 1$. To get \bar{x} we thus need an inverter that inverts each of the bits of b . We need to connect that inverter to a MUX which selects the inverted b (if subtraction) or the original b (if addition). We will add one through the use of the carry in input of the lsb since that will be 0 when addition and 1 when subtraction.

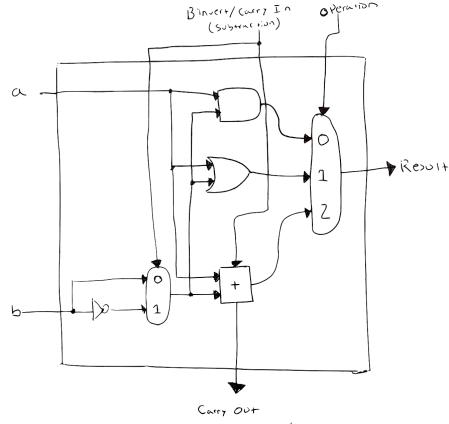


Figure 26. For subtraction we use the value of carry in to additionally select whether we should invert b or not. The operation code selects between AND, OR, and ADD (which also performs subtraction depending upon value of carry in).

5.2.2 NOR (Neither a nor b)

Because

$$(\overline{a + b}) = \bar{a} \cdot \bar{b}$$

we are able to utilize the existing Binvert and AND gate and only need to add an inversion of a .

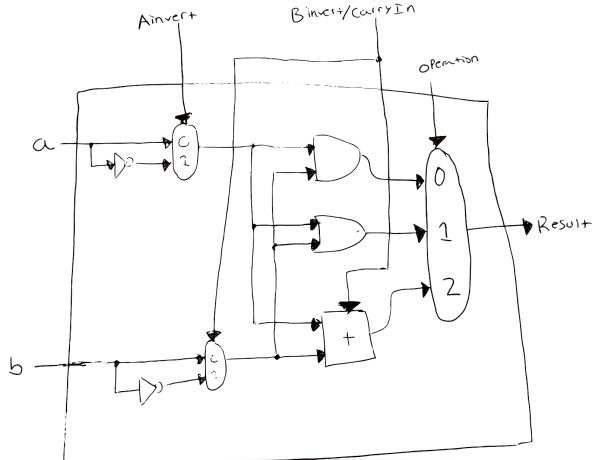


Figure 27. For NOR the output of the AND gate is used but with both Ainvert and Binvert inputs set as true.

5.2.3 Shift Less Than

For `slt` we return 1 if `rs < rt` and 0 otherwise. To test for this we can use subtraction because

$$a < b \Leftrightarrow (a - b) < 0$$

and thus if the result of $a - b$ is negative we set the lsb to 1. How does this work in practice? Well the 31 most significant ALUs are easy, we just connect 0 directly to the MUX that selects the desired operation. We will label this input `sltResult`. For the least significant bit we connect the result of the adder of the msb to the least significant bit input `sltResult`.

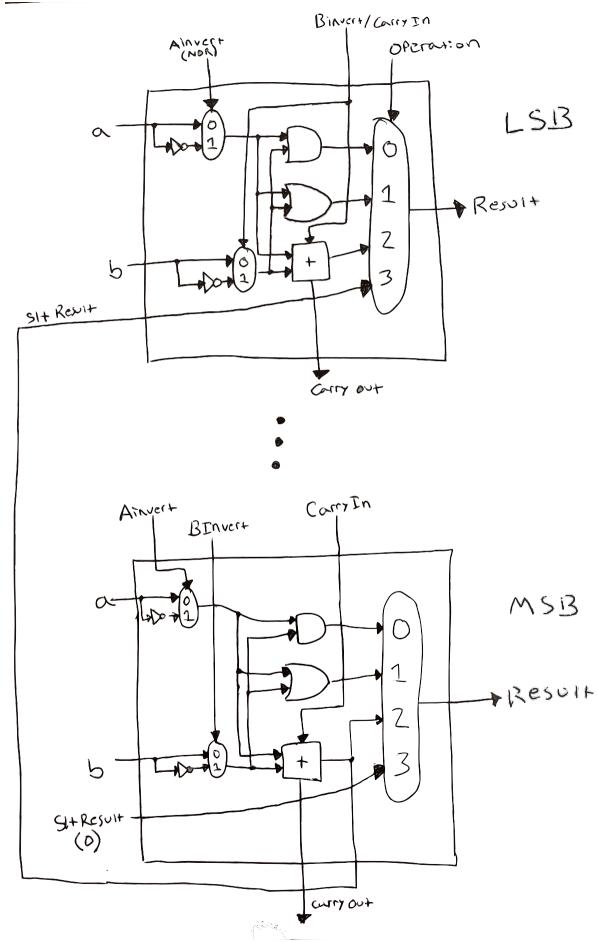


Figure 28. To perform `slt` we set operation to 3, Binvert to 1 and `sltResult` to 0 for the 31 MSBs but use the result of the adder for the MSB to set the value of `sltResult` of the LSB

SECTION 6

Memory Elements

In addition to providing combinational logic, we also want our processor to have state. Let's look at how we can create memory elements.

SUBSECTION 6.1

S-R Latch

One of the main building blocks of memory is a **pair of cross-coupled NOR gates**. Let's say you have two lights, and you wanted it so that turning one one would turn the other off and vice-versa. You could use these cross coupled NOR gates to achieve this.

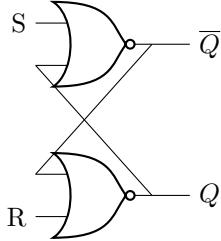


Figure 29. We have two inputs R (Reset) and S (Set). If you activate R and not S then the output of the top NOR gate will be 0 (since one of the inputs is 1). This will then connect with the bottom NOR gate and thus $\bar{Q} = 1$. The reason that the output of one gate can be the input of another is that you wait until the system is stable (since for same inputs you will always have same outputs in a combinational circuit)

SUBSECTION 6.2 Latches

6.2.1 Transparent Latch

A transparent latch is one which immediately responds to a change in the input. In this way it is an asynchronous memory element.

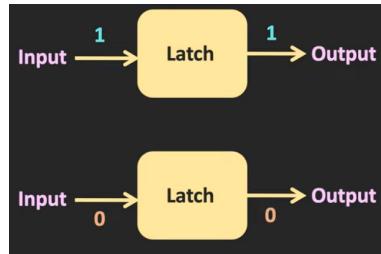


Figure 31. A transparent latch will immediately respond to a change in input. So in this example the input of a latch changes from 0 to 1 and that change is immediately reflected in the output.

Notice that the output is the same as the input. This is because the output of a latch is the same as the memory value that is stored inside.

6.2.2 Gated Latch

A gated latch is a latch whose output only changes to changes in input (becomes transparent) when the clock is high.

Inputs		Outputs	
S	R	Q	\bar{Q}
0	0	X	X
0	1	0	1
1	0	1	0
1	1	!	!

Figure 30. When both S and R are deasserted then the value of Q and \bar{Q} remain unchanged. When both are asserted, then undefined behavior is the result. When only S is asserted then $Q = 1$ and when only R is asserted then $\bar{Q} = 1$.

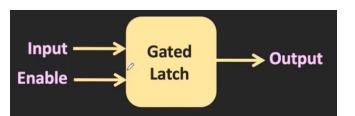


Figure 32. The gated latch has an additional “enable” input. This is asserted when the clock level is high and deasserted when not.

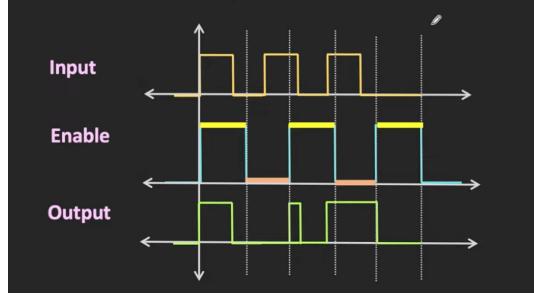


Figure 33. We see that in a gated latch that changes to input when the clock is low are not affected in the output, and that the output only changes when the clock is high.

We can utilize our S-R latch to build a gated D latch. What is our goal? We want our output Q to reflect our data input D only when our clock C is asserted. Therefore, we will have two AND gates with C as an input of each. If we connect the output of these AND gates to the input of our NOR gates then that means that if C is deasserted then the inputs of our SR Latch will be $(0, 0)$ and nothing will happen. If we connect D with the AND gate going into the set input of our SR Latch, then that will set $Q = 1$. Similarly if we have the inverse of D going into the AND gate whose output is the input for the reset, then that will turn $Q = 0$.

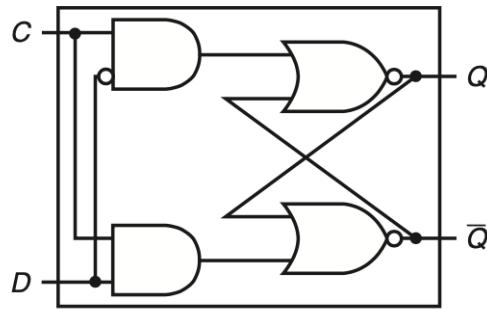


Figure 34. Implementing a D Latch using logic gates.

6.2.3 Flip-Flop

A flip-flop is very similar to a gated latch except that it becomes transparent only on the rising edge or falling edge of the clock, depending upon if it is *positive-edge triggered* or *negative-edge triggered*.

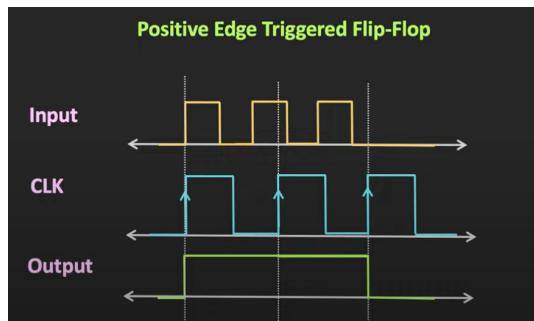


Figure 36. We see that the output only is changed when there is a rising clock edge.

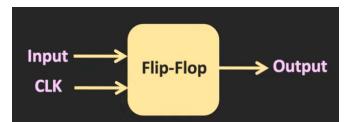
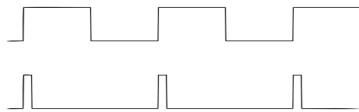


Figure 35. A flip-flop also accepts the clock signal as an input, and is transparent only on the rising or falling edge (exclusive or).

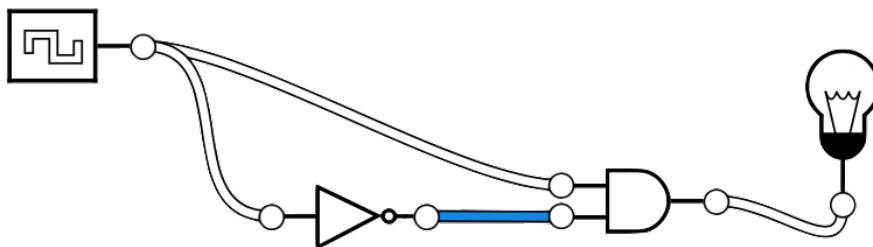
How can we implement a Flip-Flop? Well what if we had some kind of circuit that took the clock signal and converted the rising (or falling) edge into little blips:



Well, then we can use our D-Latch implementation, because it will only be transparent on rising edges (since E will only be asserted for a very short time). How can we implement this?

Making Use of Inverter Delay

Well one way is with the use of AND gate whose inputs are D and D passing through a NOT gate:



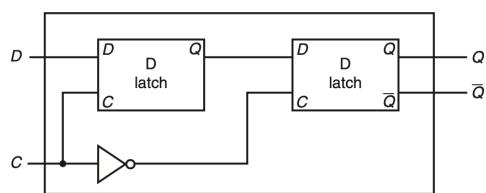
You may think that the light bulb will never turn on, and you would almost be correct. But because it takes a little bit of time for the signal to pass through the inverter, when the clock goes from low to high (rising edge), then both inputs of the and function will be 1 for a split second and the lightbulb will turn on.

Master-Slave Flip Flop

Let's say you want to design a negative-edge triggered D flip flop. Well, if we constrain ourselves to having the data input being stable for the duration of the rising and falling edge of the clock cycle (and a bit before and after), then we can create a flip flop by connecting two D-latches together:

- The first latch will be transparent on a high clock signal and take in D as the input setting Q . The output Q will then connect to another D latch.
- The second latch will be transparent on a low clock signal and take in D as the input (passed through the first D-Latch).

In this way the flip flop will only update on a falling-edge trigger.



Definition 4

D Latch: A latch with one data input (called D) that stores the value of that input signal in the internal memory when the clock is high.

D Flip-Flop: A flip-flop with one data input (called D) that stores the value of that input signal in the internal memory when the clock edge occurs.

SECTION 7

The Processor

SUBSECTION 7.1

Introduction

We will be building a simplified RISC processor that implements a subset of the RISC ISA. We will be implementing instructions related to:

- Memory-Reference: `lw` and `sw`
- ALU: `add`, `sub`, `AND`, `OR`, and `slt`
- Branch: `beq` and `j`

What does a processor look like?

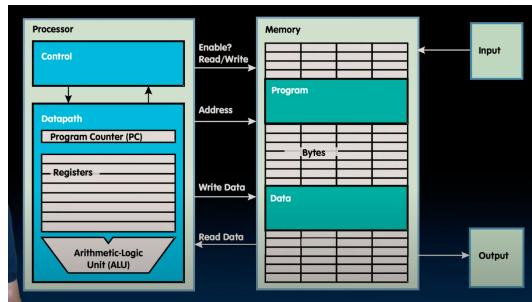


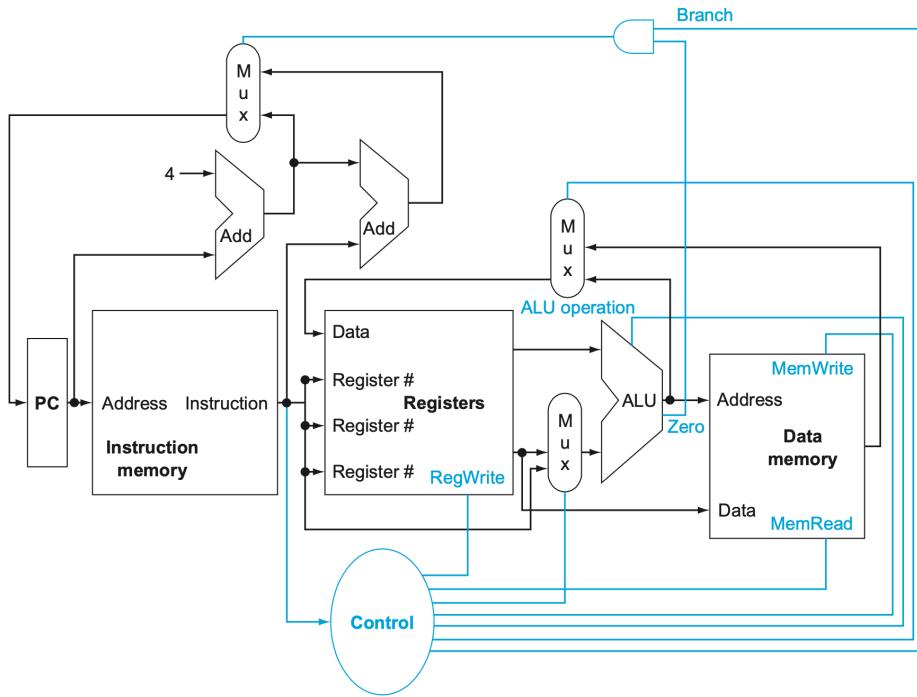
Figure 37. A big picture view of a processor

What are the two main components of a CPU:

- Datapath: This is the portion of the processor that contains the hardware necessary to perform operations required by the processor. This is the brawn of the processor. We could build a separate datapath for every instruction. Instead, we build one datapath that operates many instructions.
- Control: Another hardware portion of the processor that tells the datapath what needs to be done (the brain)

How does a processor perform our instructions? Well, let's look at a big picture view into a processor:

Processor (CPU): This is the active part of the computer that does all the work (data manipulation and decision-making)



What is going on here?

- The first thing that happens is we load the instruction using the address stored in the PC. Recall that the PC stores the address of the next instruction to be performed.
- The address from the PC also gets sent to an adder to load the next instruction.
- But if the instruction is a branch instruction then the PC needs to reflect that, so thus the MUX.
- The register fields from the instruction are then sent to the register file giving us the appropriate register operands which can then be used to perform our operation.
- The operation could be to compute a memory address (for load or store), or to perform an ALU operation.
- If it is a memory store operation, then the appropriate value is stored in memory.
- If it was an arithmetic operation or a memory load operation, the value is then written back to the appropriate register.
- Note that the ZERO output of the ALU is used on **beq** instructions.

SUBSECTION 7.2

Logic Design Conventions